

MAPR Academy

Introduction to Developing Hadoop Applications

© 2014 MapR Technologies. MAPR 1

 Learning Goals

- ▶ 1. Illustrate the MapReduce model conceptually
- 2. Describe a brief history of MapReduce
- 3. Discuss how MapReduce works at a high-level
- 4. Define how data flows in MapReduce

© 2014 MapR Technologies. MAPR 2



- Cards Game Animation Placeholder

© 2014 MapR Technologies. MAPR 3



 Learning Goals

1. Illustrate the MapReduce model conceptually
- ▶ 2. Describe a brief history of MapReduce
3. Discuss how MapReduce works at a high-level
4. Define how data flows in MapReduce

© 2014 MapR Technologies. MAPR 4

 Lisp Map-Reduce



(map square '(1 2 3 4)) = (1 4 9 16)

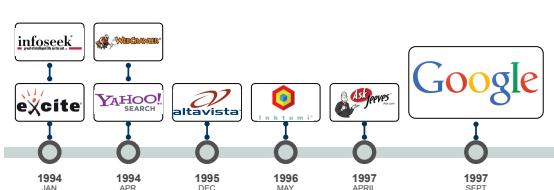
- Applies same logic to each value, *one value at a time*
- Emits result for each value

(reduce + '(1 4 9 16)) = 30

- Applies same logic to *all the values* taken together
- Emits single result for all values

© 2014 MapR Technologies. MAPR 5

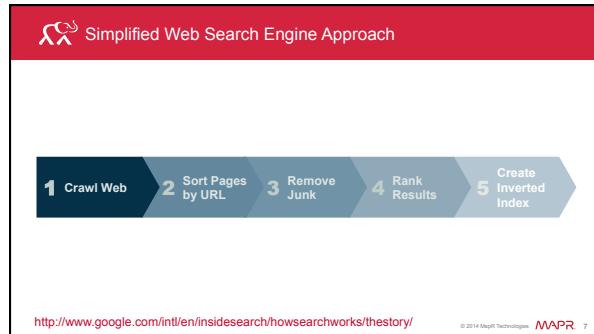
 Web Search Engines



Search Engine	Year	Month
infoseek	1994	JAN
WebCrawler	1994	APR
excite	1994	APR
YAHOO! SEARCH	1994	APR
altavista	1995	DEC
lycos	1996	MAY
Ask Jeeves	1997	APRIL
Google	1997	SEPT

© 2014 MapR Technologies. MAPR 6





 Word Count Algorithm from Google White Paper

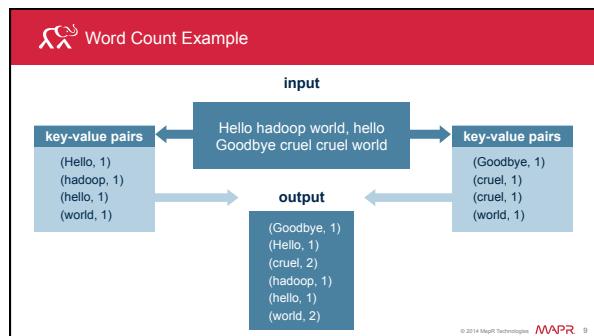
```

map(String key, String value):
    // key: document name
    // value: document contents
    for each word w in value:
        EmitIntermediate(w,"1");

reduce(String key, Iterator values):
    // key: a word
    // values: a list of counts
    int result = 0;
    for each v in values:
        result +=ParseInt(v);
    Emit(asString(result));
  
```

MapReduce: Simplified Data Processing on Large Clusters Jeffrey Dean (jeff@google.com) & Sanjay Ghemawat (sanjay@google.com), Google, Inc.

© 2014 MapR Technologies  8

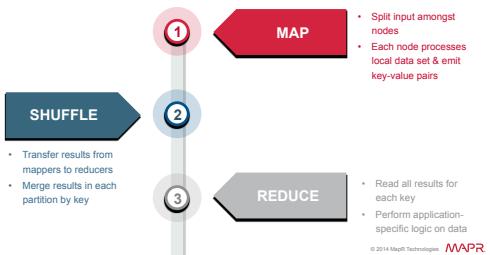


Learning Goals

1. Illustrate the MapReduce model conceptually
2. Describe a brief history of MapReduce
- ▶ 3. Discuss how MapReduce works at a high-level
4. Define how data flows in MapReduce

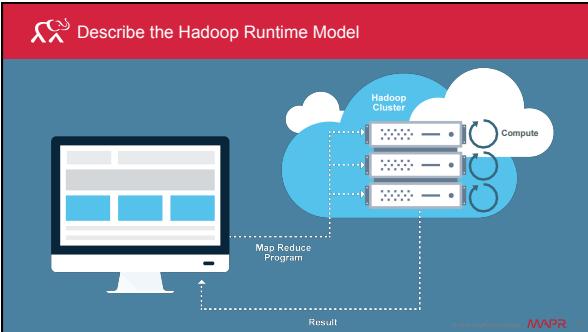
© 2014 MapR Technologies 10

Describe a Summary of Hadoop MapReduce



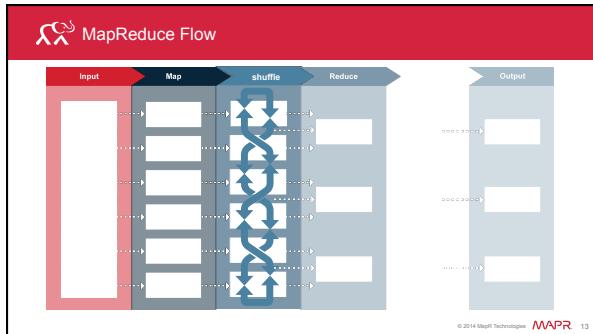
© 2014 MapR Technologies 11

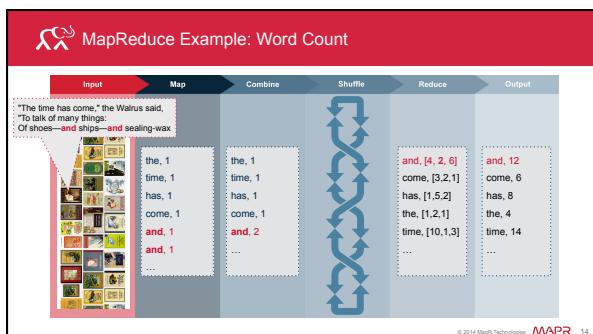
Describe the Hadoop Runtime Model



© 2014 MapR Technologies 12







Knowledge Check

The following are true of a MapReduce job:

1. The data in the map phase is split amongst the task tracker nodes where the data is located
2. The shuffle phase is handled by the Hadoop framework.
3. Output from the mappers is sent to the reducers as partitions.
4. Reducers emit zero or more key-value pairs based on the application logic
5. All of the above

© 2014 MapR Technologies MAPR 15



Learning Goals

1. Illustrate the MapReduce model conceptually
2. Describe a brief history of MapReduce
3. Discuss how MapReduce works at a high-level
- ▶ 4. Define how data flows in MapReduce

© 2014 MapR Technologies  16

Define a Simplified MapReduce Workflow

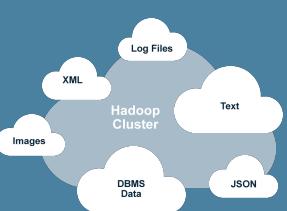
- 1 Load
- 2 Analyze
- 3 Store
- 4 Read



Describe How Data Gets Loaded Into the Cluster

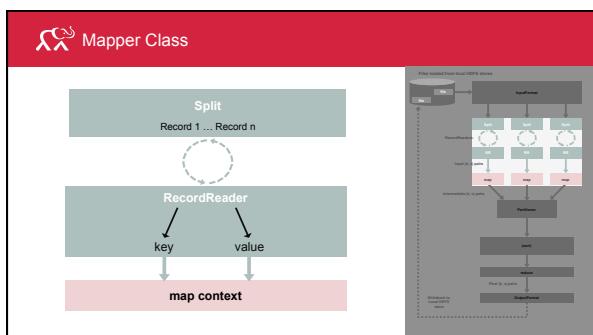
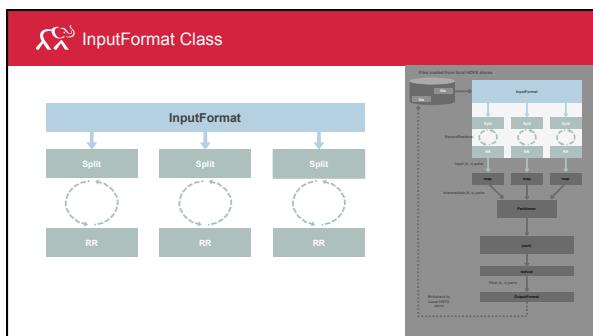
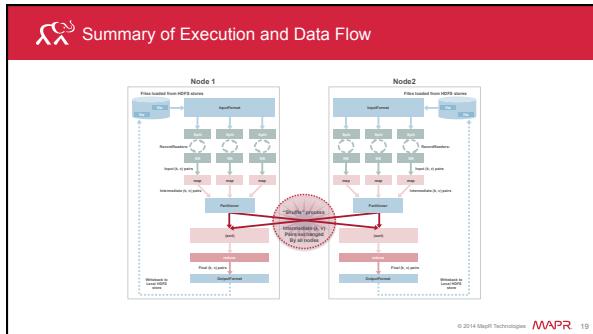
 **MapR-FS**

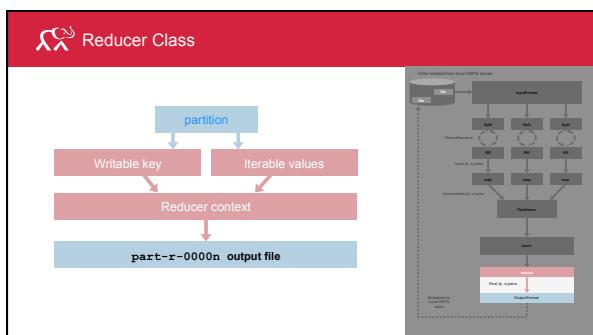
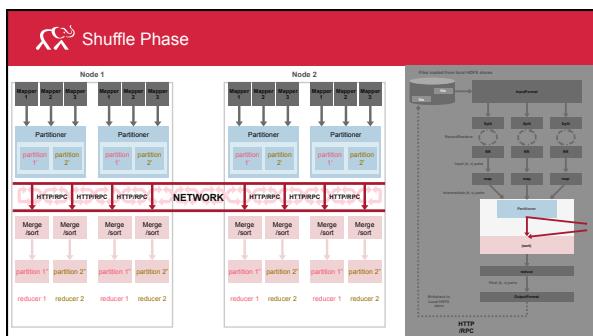
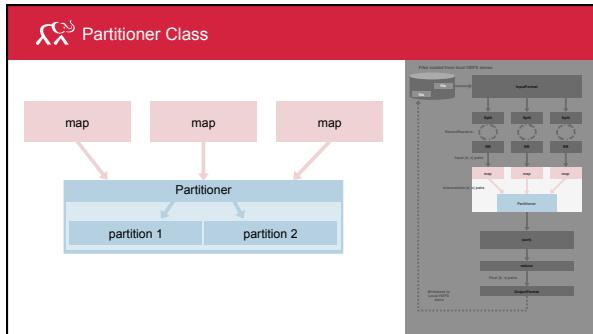
- POSIX + NFS Access
- Pre-Load or Permanent Store



© 2014 MapR Technologies  16







 Results From MapReduce Job

- _SUCCESS
- _logs/history*
- part-r-00000, part-r-00001, . . .
- part-m-00000, part-m-00001, . . .

© 2014 MapR Technologies  25

 Refer to Lab Guide for Exercise Instructions



 Summary

- MapReduce model
- History of MapReduce
- Data flow in MapReduce

© 2014 MapR Technologies  27







MapR Academy

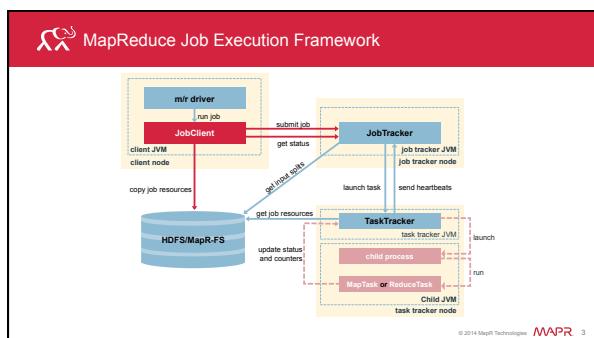
Job Execution Framework
MapReduce v1 & v2

© 2014 MapR Technologies. **MAPR** 1

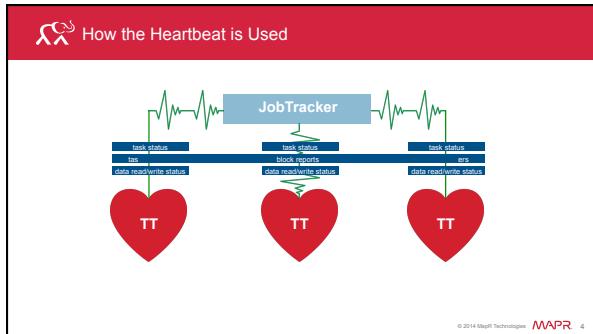
✖ Learning Goals

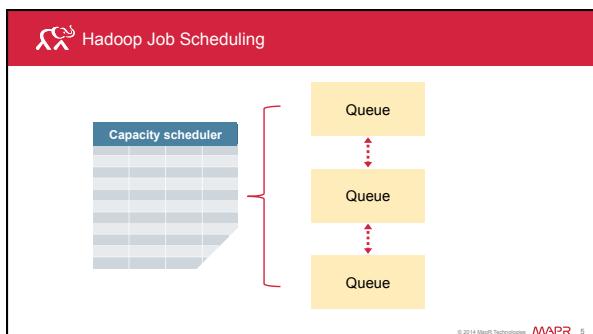
- ▶ 1. Describe the MapReduce v1 job execution framework
- 2. Compare MapReduce v1 to MapReduce v2 (YARN)
- 3. Describe how jobs execute in YARN
- 4. Describe how to monitor jobs in YARN

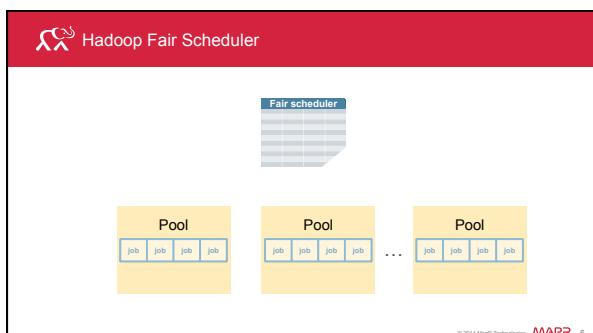
© 2014 MapR Technologies. **MAPR** 2













 Fair Scheduler Web UI

CentOS006 Fair Scheduler Administration

Pools

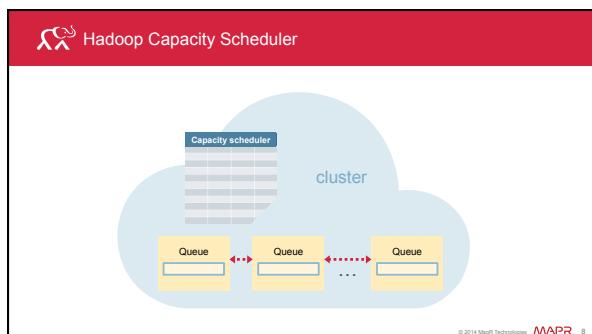
Pool	Running Jobs			Map Tasks			Reduce Tasks			Scheduling Mode
	Min Share	Max Share	Running	Fair Share	Min Share	Max Share	Running	Fair Share		
ExpressLane	0	-	0	0.0	0	-	0	0.0	FAIR	
pasalotto	1	0	-	1	0	-	0	0.0	FAIR	
c-rchar	0	0	-	0	0.0	-	0	0.0	FAIR	
user01	0	0	-	0	0.0	-	0	0.0	FAIR	
default	0	0	-	0	0.0	-	0	0.0	FAIR	

Running Jobs

Submitted	JobID	User	Name	Pool	Priority	Map Tasks	Reduce Tasks	
						Finished	Running	Fair Share
Apr 24, 11:29	job_201402251527_0018	pasalotto	pasalotto	pasalotto	NORMAL	0 / 1	1	1.0
						0 / 1	0	0.0

Scheduler runs on JobTracker node

© 2014 MapR Technologies  7



 Limitations in the Hadoop Execution Framework

Aspect				
scalability	availability	inflexibility	scheduler optimization	program support
Framework is restricted to map-reduce programs				
Framework does not optimize scheduling of jobs				
Map and reduce slots are not interchangeable				
Single job tracker and namenode introduces SPOF				
Single job tracker restricts job throughput				
Limitation				

© 2014 MapR Technologies  9



Learning Goals

1. Describe the MapReduce job execution framework
- ▶ 2. Compare MapReduce v1 to MapReduce v2 (YARN)
3. Describe how jobs execute in YARN
4. Describe how to monitor jobs in YARN

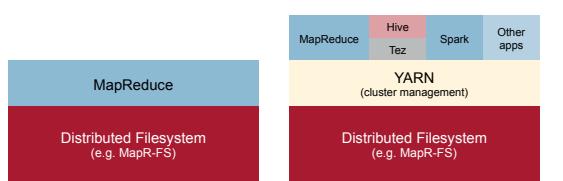
© 2014 MapR Technologies 10

Motivation for YARN

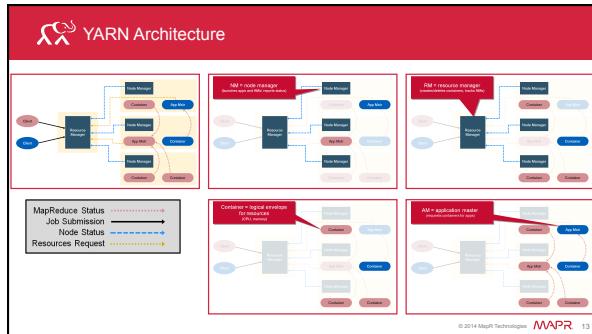
- 1 Map and reduce slot configuration is not dynamic
Inflexibility leads to underutilization
No slot configuration in YARN
- 2 Hadoop framework only supports MapReduce jobs
No support for non-MR apps
YARN supports MR and non-MR apps
- 3 Single job tracker has scalability limitations
~4000 nodes per cluster max
YARN equivalent of JT has multiple instances per cluster for scale
- 4 YARN uses same API and CLI as MRv1
(with similar Web Uis)



Differences MRv1 and MRv2



© 2014 MapR Technologies 12



Knowledge Check

The following is true of YARN

1. There is no slot configuration in YARN
2. The YARN equivalent of Job Tracker supports multiple instances per cluster to scale.
3. YARN support MapReduce and non-MapReduce jobs
- ▶ 4. All of the above

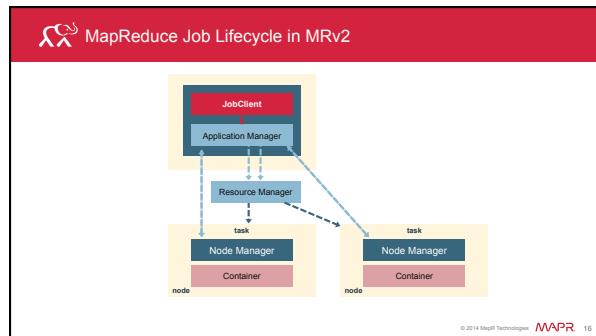
© 2014 MapR Technologies **MAPR** 14

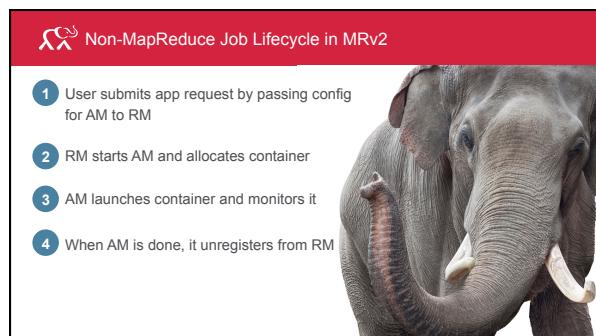
Learning Goals

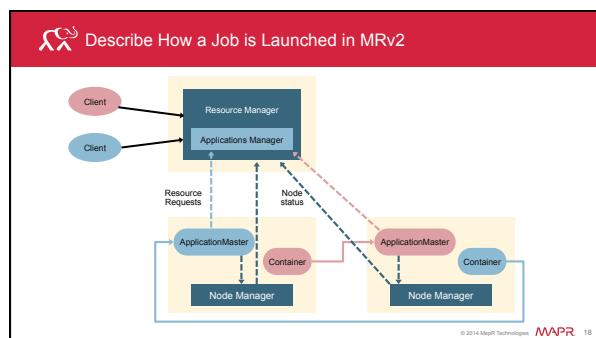
1. Describe the MapReduce job execution framework
2. Compare MapReduce v1 to MapReduce v2 (YARN)
- ▶ 3. Describe how jobs execute in YARN
4. Describe how to monitor jobs in YARN

© 2014 MapR Technologies **MAPR** 15

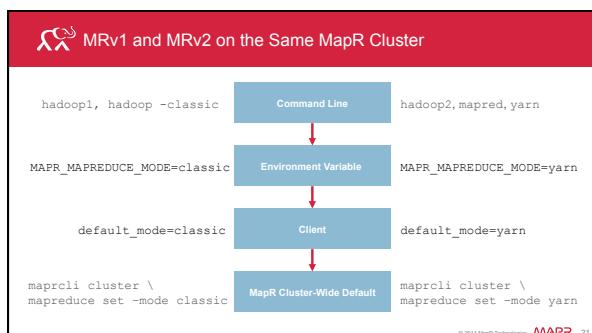
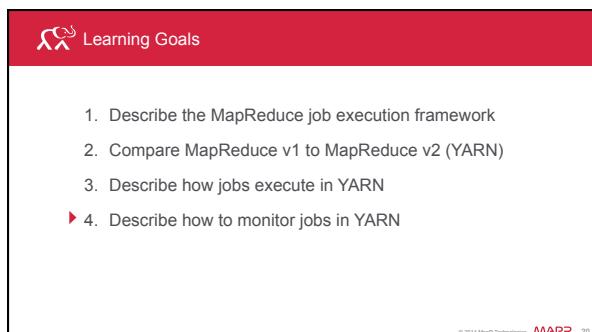
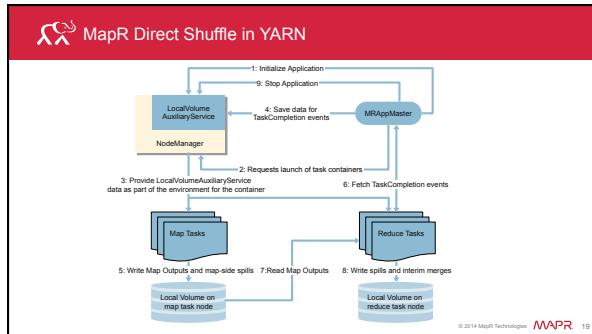


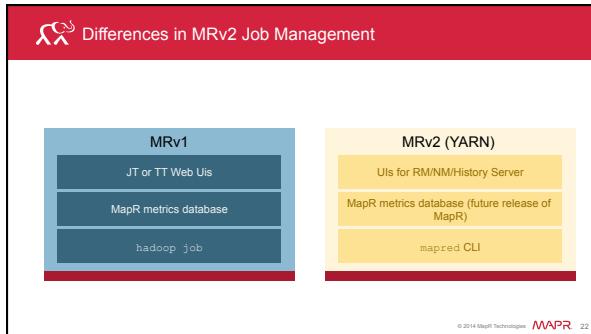


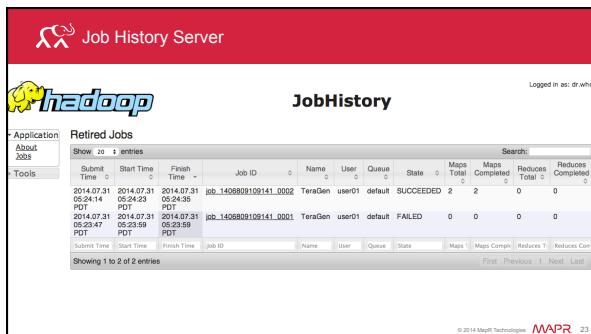










 Job History Server

 **JobHistory**

Application
Jobs
2008
Tools

Retired Jobs

Submit Time	Start Time	Finish Time	Job ID	Name	User	Queue	State	Maps Total	Maps Completed	Reduces Total	Reduces Completed
2014-07-31 05:24:14 PDT	2014-07-31 05:24:33 PDT	2014-07-31 05:24:38 PDT	pb_1406809109141_0002	TeraGen	user01	default	SUCCEEDED	2	2	0	0
2014-07-31 05:23:47 PDT	2014-07-31 05:23:59 PDT	2014-07-31 05:24:07 PDT	pb_1406809109141_0001	TeraGen	user01	default	FAILED	0	0	0	0

Show: entries Search:

Submit Time | Start Time | Finish Time | Job ID | Name | User | Queue | State | Maps Total | Maps Completed | Reduces Total | Reduces Completed

Showing 1 to 2 of 2 entries

© 2014 MapR Technologies  23

 Job History (2)

 **MapReduce Job**
job_1406809109141_0002

Application
Jobs
Over-size
Counters
Configuration
Map Tasks
Reduce Tasks
Tools

Job Overview

Job Name: TeraGen
User: user01
Queue: default
Status: SUCCEEDED
Uberized: false
Submitted: Thu Jul 31 05:24:14 PDT 2014
Started: Thu Jul 31 05:24:23 PDT 2014
Finished: Thu Jul 31 05:24:38 PDT 2014
Elapsed: 11sec
Diagnostics:
Average Map Time: 8ms

ApplicationMaster

Attempt Number	Start Time	Node	Logs
1	Thu Jul 31 05:24:19 PDT 2014	mapr1node:8042	logs

Total Task Type

Map	Reduce	Total	Complete
2	0	2	0
Attempts			
Map	Failed	0	Killed
Reduces	Successful	0	0

24



Job History (3)

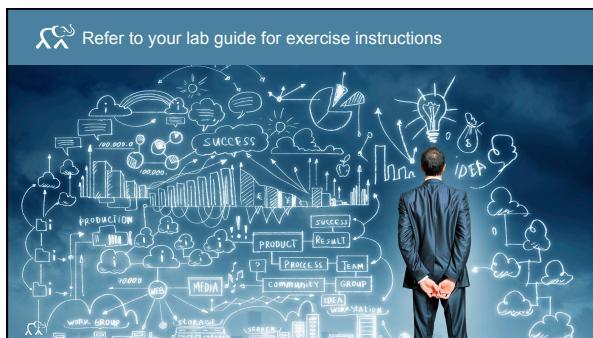
Counters for job_1406809109141_0002

Counter Group	Name	Counters
File System Counters	FILE: Number of bytes read	Map: 0, Reduce: 0, Total: 0
	FILE: Number of bytes written	Map: 130398, Reduce: 0, Total: 130398
	FILE: Number of large read operations	Map: 0, Reduce: 0, Total: 0
	FILE: Number of large write operations	Map: 0, Reduce: 0, Total: 0
	FILE: Number of small read operations	Map: 0, Reduce: 0, Total: 0
	MAPRED: Number of bytes read	Map: 164, Reduce: 164, Total: 164
	MAPRED: Number of bytes written	Map: 100000, Reduce: 0, Total: 100000
	MAPRED: Number of large read operations	Map: 0, Reduce: 0, Total: 0
	MAPRED: Number of large write operations	Map: 14, Reduce: 14, Total: 14
	MAPRED: Number of small read operations	Map: 2012, Reduce: 0, Total: 2012
Job Counters	Launched map tasks	Map: 0, Reduce: 0, Total: 0
	Other local map tasks	Map: 0, Reduce: 0, Total: 0
	Total time spent by all maps in occupied slots (ms)	Map: 0, Reduce: 0, Total: 0
	CPU time spent (ms)	Map: 320, Reduce: 0, Total: 320
	File output bytes	Map: 0, Reduce: 0, Total: 0
Map-Reduce Framework	GC time elapsed (ms)	Map: 56, Reduce: 56, Total: 56
	Input split bytes	Map: 164, Reduce: 164, Total: 164
	Map output records	Map: 1000, Reduce: 0, Total: 1000

Job History (4)

Configuration for MapReduce Job job_1406809109141_0002

key	value	source chain
dfs.blocksize	268435456	job.xml → default.xml → core-default.xml
dfs.bytes-per-checksum	512	job.xml → default.xml → core-default.xml
dfs.ha.fencing.ssh.connect.timeout	30000	job.xml → default.xml → core-default.xml
dfs.namenode.checkpoint.dir	\$maprfs:/tmp/staging/history/done/2014/07/31/000000/job_1406809109141_0002/conf.xml	job.xml → default.xml → core-default.xml
dfs.namenode.checkpoint.edits.dir	\$maprfs:/tmp/staging/history/done/2014/07/31/000000/job_1406809109141_0002/conf.xml	job.xml → default.xml → core-default.xml
dfs.namenode.checkpoint.period	3600	job.xml → default.xml → core-default.xml
file.blocksize	67108864	job.xml → default.xml → core-default.xml
file.bytes-per-checksum	512	job.xml → default.xml → core-default.xml
file.client-write-packet-size	65536	job.xml → default.xml → core-default.xml
file.replication	1	job.xml → default.xml → core-default.xml



Summary

- Execution of MapReduce
- Motivation of YARN development
- Differences between the 2 versions
- YARN job execution and management

© 2014 MapR Technologies **MAPR** 28

Next Steps



© 2014 MapR Technologies **MAPR** 29



MAPR Academy

Write a MapReduce Program

© 2014 MapR Technologies MAPR 1

Learning Goals

- ▶ 1. Summarize programming problem
- 2. Design and implement
 - Mapper class
 - Reducer class
 - Driver class
- 3. Build and execute code

© 2014 MapR Technologies MAPR 2

MapReduce Design Patterns

Pattern			
Summarizing data	Filtering data	Organizing data	Joining data
Map-side and reduce-side joins			
Transform, partition, sort, and generate data			
Sample, sanitize, identify top n, and filter unique data			
Statistical summaries, counts, & indexes for groups of data			
Notes			

© 2014 MapR Technologies MAPR 3



Some MapReduce Programming Tips

- Start with a template for the driver, mapper, and reducer classes
- Modify the template to suit the needs of your application
- Understand the flow and transformation of data:**



- Identify appropriate types for keys and values

© 2014 MapR Technologies 4

Example Data Set

2000	2025191	1788950	236241	1544607	1458185	86422	480584	330765	149819
2001	1991082	1862846	128236	1483563	1516008	-32445	507519	346638	160681
2002	1853136	2010894	-157758	1337815	1655232	-317417	515321	356662	159659
2003	1782314	2159899	-377585	1258472	1796890	-538418	523842	363009	160833
2004	1880114	2292841	-412727	1345369	1913330	-567961	534745	379511	155234
2005	2153611	2471957	-318346	1576135	2069746	-493611	577476	402211	175265
2006	2406869	2655050	-24818	1798487	2232981	-434494	608382	422069	186313
2007	2567985	2728686	-160701	1932898	2275049	-342153	635089	453637	181452
2008	2523991	2982544	-458553	1865945	2507793	-641848	658046	474751	183295
2009	2104989	3517677	-1412688	1450994	3000661	-1549681	654009	517016	136993
2010	2162706	3457079	-1294373	1531019	2902397	-1371378	631687	554682	77005
2011	2303466	3603059	-1299593	1737678	3104453	-1366775	565788	498606	67182
2012	2450164	3537127	-1086963	1880663	3029539	-1148876	569501	507588	61913

© 2014 MapR Technologies 5

Learning Goals

- Summarize programming problem
- Design and implement
 - Mapper class
 - Reducer class
 - Driver class
- Build and execute code

© 2014 MapR Technologies 6



 Input to the Mapper Class

Input format = TextInputFormat
Key = LongWritable
Value = Text

More detail later

First record:

input key input value

0 1901 588 525 63 588 525 63

```
 StringTokenizer itr = new StringTokenizer(value.toString(),"\s+");
```

© 2014 MapR Technologies  7

 Output of the Mapper Class

output key output value

```
context.write(new Text("summary"),new Text(year + " - " + delta));
```

First few lines of output

summary 1901_63
summary 1902_77
summary 1903_45
summary 1904_-43
summary 1905_-23
summary 1906_25

© 2014 MapR Technologies  8

 Design and Implement the Mapper Class

```
public class ReceiptsMapper extends Mapper<LongWritable,Text,Text> {
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        StringTokenizer iterator = new StringTokenizer(value.toString(),"\s+");
        String year = iterator.nextToken();
        iterator.nextToken();
        iterator.nextToken();
        String delta= iterator.nextToken();
        context.write(new Text("summary"), new Text(year + " - " + delta));
    }
}
```

© 2014 MapR Technologies  9





Knowledge Check

Which statements are true of Mapper class?

1. The four arguments to the Mapper class represent input key type, input value type, output key type, and output value type.
2. The mapper class calls the map() method
3. The first two arguments to the map() method are the key and value which must match the types defined in the Mapper class definition
- ▶ 4. All of the above

© 2014 MapR Technologies. 10



Learning Goals

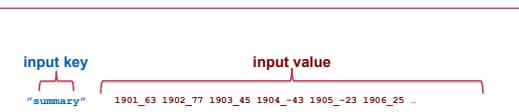
1. Summarize programming problem
2. Design and implement
 - Mapper class
 - ▶ – Reducer class
 - Driver class
3. Build and execute code

© 2014 MapR Technologies. 11



Input to the Reducer Class

- Mapper output key = Text → Reducer input key = Text
- Mapper output value = Text → Reducer input values = Text


© 2014 MapR Technologies. 12



 Output of the Reducer Class

```

Output key = Text
• min(year):
    |
    +--> min(2009)

Output value = FloatWritable
min (from long to FloatWritable)
    |
    +--> -1412688.0

```

© 2014 MapR Technologies  13

 Design and Implement the Reducer Class (1)

```

public class ReceiptsReducer extends Reducer
<Text,Text,Text,FloatWritable> {
    public void reduce(Text key, Iterable<Text> values, Context context)
        throws IOException, InterruptedException {
        long tempValue = 0L, minValue=Long.MAX_VALUE;
        Text tempYear=null, tempValue=null, minYear=null, maxYear=null;
        String compositeString;
        String[] compositeStringArray;
    }
}

```

© 2014 MapR Technologies  14

 Design and Implement Reducer Class (2)

```

for (Text value: values) {
    compositeString = value.toString();
    compositeStringArray = compositeString.split(" ");
    tempYear = new Text(compositeStringArray[0]);
    tempValue = new Long(compositeStringArray[1]).longValue();
    if(tempValue < min) {
        min=tempValue;
        minYear=tempYear;
    }
}
Text keyText = new Text("min" + "(" + minYear.toString() + "): ");
context.write(keyText, new FloatWritable(min));
}
}

```

© 2014 MapR Technologies  15





Knowledge Check

- Which statements are true of the Reducer class?
1. The four arguments to the Reducer class represent the type for input key, input value, output key, and output value
 2. If the output value of the Mapper is Text, the input value to the Reducer can be LongWritable
 3. If the output key of the Mapper is Text, the input key to the Reducer class must also be Text.
 4. All of the above
 - ▶ 5. 1 & 3 only

© 2014 MapR Technologies. 16



Learning Goals

1. Summarize programming problem
2. Design and implement
 - Mapper class
 - Reducer class
 - ▶ - Driver class
3. Build and execute code

© 2014 MapR Technologies. 17



Implement the Driver (1)

```
public class ReceiptsDriver extends Configured implements Tool {
    public int run(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.printf("usage: %s [generic options] <inputfile>
<outputdir>\n", getClass().getSimpleName());
            System.exit(1);
        }
        Job job = new Job(getConf(), "my receipts");
        job.setJarByClass(ReceiptsDriver.class);
        job.setMapperClass(ReceiptsMapper.class);
        job.setReducerClass(ReceiptsReducer.class);
    }
}
```

© 2014 MapR Technologies. 18



Implement the Driver (2)

```

job.setInputFormatClass(TextInputFormat.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(FloatWritable.class);
job.setMapOutputValueClass(Text.class);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
return job.waitForCompletion(true) ? 0 : 1;
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    System.exit(ToolRunner.run(conf, new ReceiptsDriver(), args));
}
}

```

© 2014 MapR Technologies 19

Knowledge Check

Which statements are true of the Driver class?

1. The Driver class first checks the invocation of the command (checks the count of the command-line arguments provided)
2. It sets values for the job, including the driver, mapper, and reducer classes used.
3. In the Driver class, we can specify how we want to launch the job – synchronously or asynchronously
- ▶ 4. All of the above

© 2014 MapR Technologies 20

Learning Goals

1. Summarize programming problem
2. Design and implement
 - Mapper class
 - Reducer class
 - Driver class
- ▶ 3. Build and execute code

© 2014 MapR Technologies 21



Configure the Environment

```
$ cat ~/.profile
export HADOOP_HOME=/opt/mapr/hadoop/hadoop-0.20.2
export LD_LIBRARY_PATH=$HADOOP_HOME/lib/native/Linux-amd64-64
export PATH=$HADOOP_HOME/bin:$PATH
export CLASSPATH=$HADOOP_HOME/*:$HADOOP_HOME/lib/*
export HADOOP_CLASSPATH=$CLASSPATH
```

© 2014 MapR Technologies  22

Build the Jar

```
$ mkdir classes
$ javac -d classes ReceiptsMapper.java
$ javac -d classes ReceiptsReducer.java
$ jar -cvf Receipts.jar -C classes/
$ javac -classpath $CLASSPATH:Receipts.jar -d classes ReceiptsDriver.java
$ jar -uvf Receipts.jar -C classes/ .
```

© 2014 MapR Technologies  23

Launch the Hadoop Job

```
$ hadoop jar Receipts.jar Receipts.ReceiptsDriver \
/user/user01/RECEIPTS/DATA/receipts.txt \
/user/user01/RECEIPTS/OUT
```

© 2014 MapR Technologies  24



Examine the Output

```
$ hadoop fs -cat /user/user01/RECEIPTS/OUT/part-r-00000  
min(2009): -1412688.0
```

© 2014 MapR Technologies **MAPR** 25

Refer to Lab Guide for Exercise Instructions



Summary

- Design and implement a Mapper, Reducer and Driver classes
- Build and execute code

© 2014 MapR Technologies **MAPR** 27







MAPR® Academy

Use the MapReduce API

© 2014 MapR Technologies. MAPR. 1

 Learning Goals

- ▶ 1. API overview
- 2. Examine
 - Mapper input processing data flow
 - Reducer output processing data flow
- 3. Explore the Mapper, Reducer and Job class API

© 2014 MapR Technologies. MAPR. 2

 Hadoop Version Support on MapR



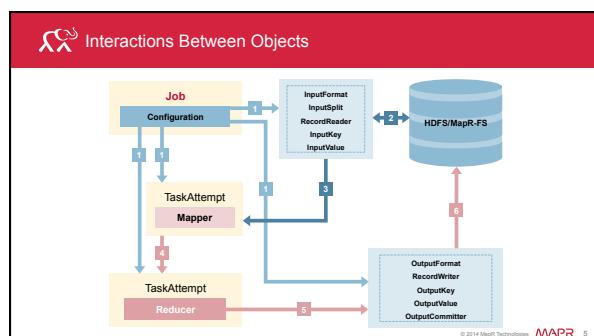
© 2014 MapR Technologies. MAPR. 3



mapred and mapreduce

Aspect	mapred	mapreduce
Supported on MapR	yes	yes
Deprecated	no	no
YARN-compatible	yes	yes
Types	interfaces	abstract classes
Objects	OutputCollector Reporter JobConf	Context
Methods	map() reduce()	map() reduce() cleanup() setup() run()
Output files	part-nnnnn	part-r-nnnnn part-m-nnnnn
Reducer input values	java.lang.Iterable	java.lang.Iterator

© 2014 MapR Technologies 4



WritableComparable Types

```

public interface WritableComparable extends Writable, Comparable
{
    void readFields(DataInput in);
    void write(DataOutput out);
    int compareTo(WritableComparable o);
}

```

© 2014 MapR Technologies 6

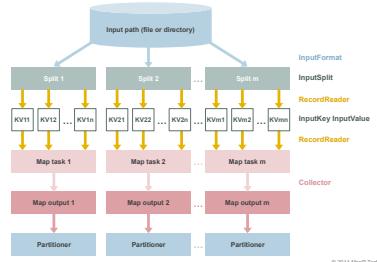


Learning Goals

1. API overview
2. Examine
 - ▶ - Mapper input processing data flow
 - Reducer output processing data flow
3. Explore the Mapper, Reducer and Job class API

© 2014 MapR Technologies 7

Mapper Input Flow



© 2014 MapR Technologies 8

InputFormat Class

```

public abstract class InputFormat<K, V> {

    // returns array of InputSplits for job
    public abstract List<InputSplit> getSplits(JobContext) throws
    IOException, InterruptedException;

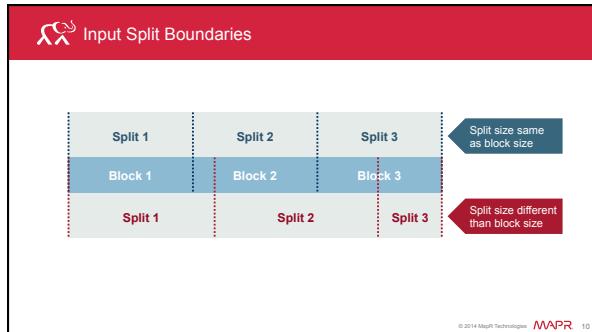
    // create new record reader
    public abstract RecordReader<K, V> createRecordReader(InputSplit
    split, TaskAttemptContext context) throws IOException,
    InterruptedException;
}
  
```

Implementations include:

- TextInputFormat (for single-line records in text files)
- SequenceFileInputFormat (for binary files)

© 2014 MapR Technologies 9





 InputSplit Class

```
public abstract class InputSplit {
    // returns number of bytes in the split
    public abstract long getLength() throws IOException;

    // returns array of nodes
    public abstract String[] getLocations() throws IOException;
}
```

Implementations include:

- FileSplit (breaks each file into splits)
- CombineFileSplit (breaks multiple files into single split)

© 2014 MapR Technologies  11

 RecordReader Interface

```
public interface RecordReader<K, V> {
    //Reads the next key/value pair from the input for processing.
    boolean next(K key, V value) throws IOException;
    //Creates an object of the appropriate type to be used as a key and value.
    K createKey();
    V createValue();
    //returns the current position in the input
    long getPos() throws IOException;
    //Close this InputSplit
    public void close() throws IOException;
    float getProgress() throws IOException;
}
```

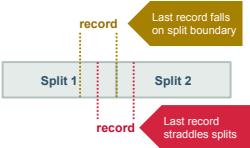
Implementations include:

- LineRecordReader
- SequenceFileRecordReader

© 2014 MapR Technologies  12



Identifying Record Boundaries



© 2014 MapR Technologies. **MAPR**. 13

Knowledge Check

In the Mapper Input flow (pick all that apply)

1. Input files are split and key-value pairs generated for each split and sent to the same Mapper.
2. The map() method is called for each key-value pair once and output sent to the partitioner.
3. Results from the mapper are stored in the file system of the Control node
4. 1, 2 & 3
- 5. 1 & 2 only

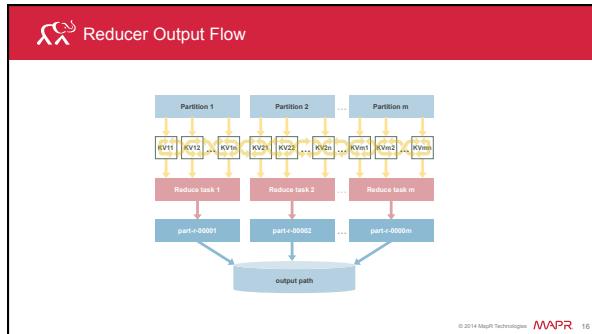
© 2014 MapR Technologies. **MAPR**. 14

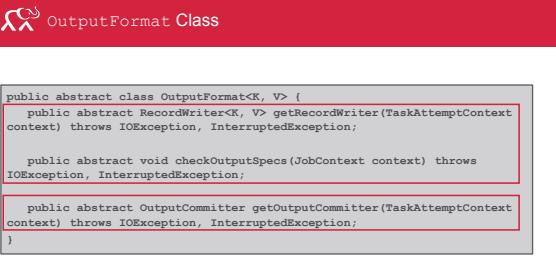
Learning Goals

1. API overview
2. Examine
 - Mapper input processing data flow
 - - Reducer output processing data flow
3. Explore the Mapper, Reducer and Job class API

© 2014 MapR Technologies. **MAPR**. 15





 OutputFormat Class

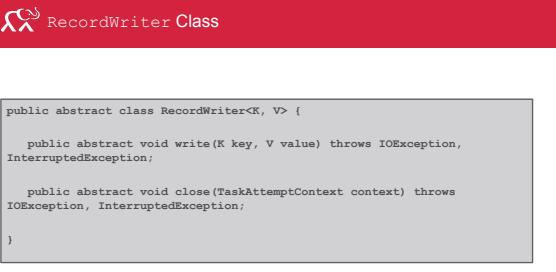
```

public abstract class OutputFormat<K, V> {
    public abstract RecordWriter<K, V> getRecordWriter(TaskAttemptContext context) throws IOException, InterruptedException;

    public abstract void checkOutputSpecs(JobContext context) throws IOException, InterruptedException;

    public abstract OutputCommitter getOutputCommitter(TaskAttemptContext context) throws IOException, InterruptedException;
}
  
```

© 2014 MapR Technologies  17

 RecordWriter Class

```

public abstract class RecordWriter<K, V> {
    public abstract void write(K key, V value) throws IOException, InterruptedException;

    public abstract void close(TaskAttemptContext context) throws IOException, InterruptedException;
}
  
```

© 2014 MapR Technologies  18



Knowledge Check

In the Reduce output flow (pick all that apply)

1. Each reducer takes its partition as input, processes one iterable list of key-value pairs at a time.
2. Produces an output file called 'part-r-0000x'
3. The output directory for this file is specified in the Job configuration and must reside on the HDFS/Mapr-FS file system
- ▶ 4. 1, 2 & 3
5. 1 & 3 only

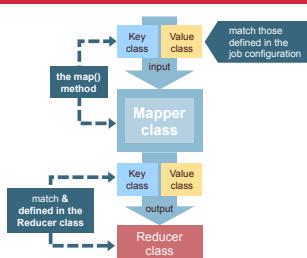
© 2014 MapR Technologies 19

Learning Goals

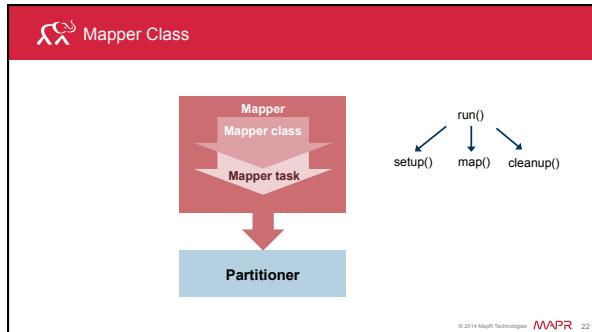
1. API overview
2. Examine
 - Mapper input processing data flow
 - Reducer output processing data flow
- ▶ 3. Explore the Mapper, Reducer and Job class API

© 2014 MapR Technologies 20

Input and Output Keys and Values


© 2014 MapR Technologies 21





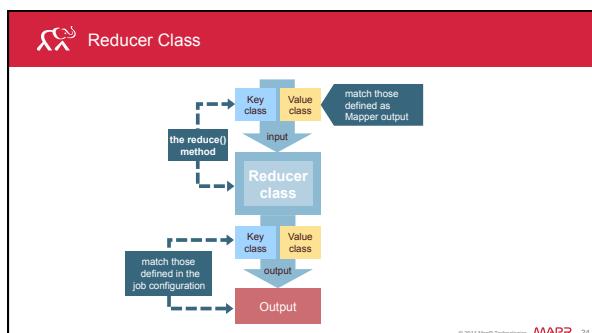
Mapper Example

```

public class MyWordcountMapper extends Mapper<LongWritable, Text, Text, IntWritable> {
    private Text word = new Text();
    private final static IntWritable one = new IntWritable(1);
    public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
        String line = value.toString();
        StringTokenizer iterator = new StringTokenizer(line);
        while (iterator.hasMoreTokens()) {
            word.set(iterator.nextToken());
            context.write(word, one);
        }
    }
}

```

© 2014 MapR Technologies 23



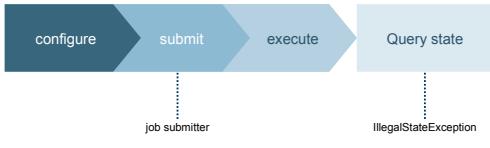


Reducer Example

```
public class MyWordCountReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable value : values) {
            sum += value.get();
        }
        context.write(key, new IntWritable(sum));
    }
}
```

© 2014 MapR Technologies  25

Job Class



```

graph LR
    A[configure] --> B[submit]
    B --> C[execute]
    C --> D[Query state]
    D -.-> E[IllegalStateException]
    E -.-> C
    
```

The flowchart illustrates the life cycle of a job. It starts with the 'configure' step, followed by 'submit', then 'execute'. After 'execute', the job enters a 'Query state' phase. From 'Query state', an arrow points down to 'IllegalStateException', which then loops back to the 'execute' step.

job submitter

IllegalStateException

© 2014 MapR Technologies  26

Job Object

```
Configuration conf = new Configuration();
Job job = new Job(conf, "mywordcount");
```

© 2014 MapR Technologies  27



Job Methods

```

void setMapperClass<T extends Mapper>(Class<T> mapper)
    Define the comparator function which key is mapped to for a single call to
    Reducer.reducer(Map<key, Iterable<Object>, Mapper, Reducer, Reducer.Context>)
void setInputFormatClass(Class<? extends InputFormat> cls)
    Set the InputFormat for the job.
void setJarPath(String path)
    Set the Jar by finding where a given class came from.
void setJobName(String name)
    Set the user-specified Job name.
void setKeyClass<T extends Writable>(Class<T> theClass)
    Set the key class for the map output data.
void setValueClass<T extends Writable>(Class<T> theClass)
    Set the value class for the map output data.
void setMapperClass<T extends Mapper>(Mapper mapper)
    Set the Mapper for the job.
void setMapReduceTasks(int tasks)
    Set the number of mappers for the job.
void setOutputFormatClass(Class<? extends OutputFormat> cls)
    Set the OutputFormat for the job.
void setMapperClass<T extends Mapper>(theClass)
    Set the key class for the job output data.
void setValueClass<T extends Writable>(theClass)
    Set the value class for job output.

```

© 2014 MapR Technologies **MAPR** 28

Job Methods

```

void setPartitionerClass(Class<? extends Partitioner> cls)
    Set the Partitioner for the job.
void setReducerClass(Class<? extends Reducer> cls)
    Set the Reducer for the job.
void setKeyComparatorClass(Class<? extends KeyComparator> cls)
    Define the comparator that controls how the keys are sorted before they are passed to the Reducer.
void setMapDir(String Path dir)
    Set the current working directory for the default file system.
void submit()
    Submit the job to the cluster and return immediately.
boolean waitForCompletion(boolean verbose)
    Submit the job to the cluster and wait for it to finish.

```

© 2014 MapR Technologies **MAPR** 29

Simplistic Driver Implementation

```

public class MyWordcountDriver {
    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = new Job(conf, "mywordcount");
        job.setJarByClass(MyWordcountDriver.class);
        job.setMapperClass(MyWordcountMapper.class);
        job.setReducerClass(MyWordcountReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

```



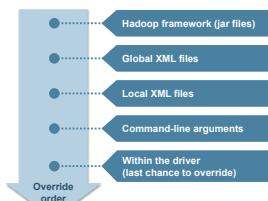
Best Way to Implement the Driver

```
public class MyWordcountDriver extends Configured implements Tool {
    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        System.exit(ToolRunner.run(conf, new MyWordcountDriver (), args));
    }

    public int run(String[] args) throws Exception {
        Job job = new Job(conf, "mywordcount");
        ...
        job.waitForCompletion(true) ? 0 : 1;
    }
}
```

© 2014 MapR Technologies 31

How to Set Configuration Parameters



More information for user configurable parameters at
<http://doc.mapr.com/display/MapR/MapR+Parameters>

© 2014 MapR Technologies 32

Refer to your lab guide for exercise instructions



Summary

- MapReduce API
- Writable types
- Mapper - Split
- Map method – single record
- Reducer - one or more partitions
- Reduce method – single key (and all associated values)
- ToolRunner for driver code

© 2014 MapR Technologies. MAPR. 34

Next Steps



© 2014 MapR Technologies. MAPR. 35

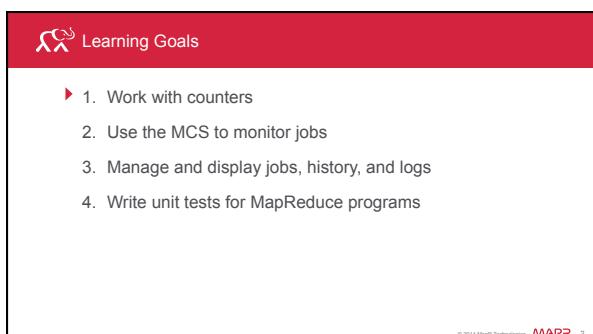




MAPR® Academy

Managing, Monitoring, and Testing MapReduce Jobs

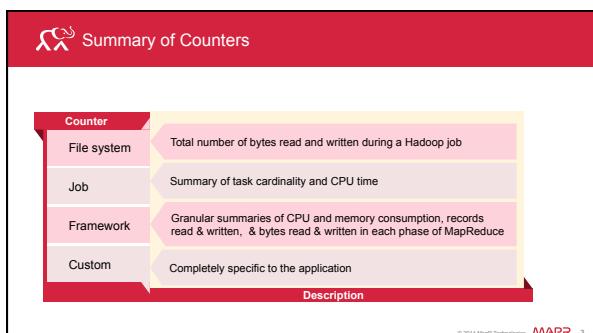
© 2014 MapR Technologies. **MAPR**. 1



 **Learning Goals**

- ▶ 1. Work with counters
- 2. Use the MCS to monitor jobs
- 3. Manage and display jobs, history, and logs
- 4. Write unit tests for MapReduce programs

© 2014 MapR Technologies. **MAPR**. 2



 **Summary of Counters**

Counter	Description
File system	Total number of bytes read and written during a Hadoop job
Job	Summary of task cardinality and CPU time
Framework	Granular summaries of CPU and memory consumption, records read & written, & bytes read & written in each phase of MapReduce
Custom	Completely specific to the application

© 2014 MapR Technologies. **MAPR**. 3



 File System Counters

Counter	Description
FILE_BYTES_WRITTEN	Total number of bytes written to local file system
MAPRFS_BYTES_READ	Total number of bytes read from MapR-FS
MAPRFS_BYTES_WRITTEN	Total number of bytes written to MapR-FS

© 2014 MapR Technologies  4

 Job Counters

Counter	Description
DATA_LOCAL_MAPS	Total number of map tasks executed on local data
FALLOW_SLOTS_MILLIS_MAPS	Total time map tasks spend waiting after slots are reserved (pre-emption)
FALLOW_SLOTS_MILLIS_REDUCES	Total time reduce tasks spend waiting after slots are reserved (pre-emption)
SLOTS_MILLIS_MAPS	Total time map tasks spend executing
SLOTS_MILLIS_REDUCES	Total time reduce tasks spend executing
TOTAL_LAUNCHED_MAPS	Total number of map tasks launched, including failed tasks
TOTAL_LAUNCHED_REDUCES	Total number of reduce tasks launched, including failed tasks

© 2014 MapR Technologies  5

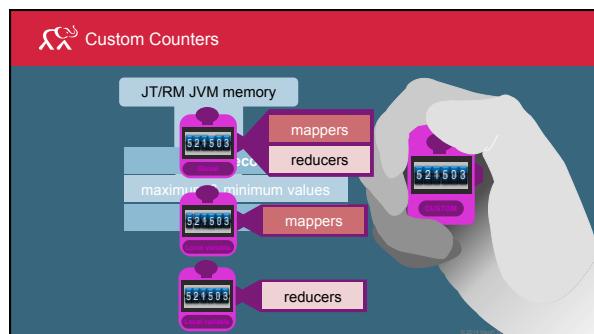
 Framework Counters (1)

Counter	Description
COMBINE_INPUT_RECORDS	Incremented for every record read during combine phase, if used.
COMBINE_OUTPUT_RECORDS	Incremented for every record written during combine phase, if used.
CPU_MILLISECONDS	Total time spent by all tasks on CPU
GC_TIME_MILLIS	Total time spent doing garbage collection
MAP_INPUT_RECORDS	Incremented for every record successfully read in map phase
MAP_OUTPUT_RECORDS	Incremented for every record successfully written in map phase
PHYSICAL_MEMORY_BYTES	Total physical memory consumed by all tasks

© 2014 MapR Technologies  6



Framework Counters (2)	
Counter	Description
REDUCE_INPUT_GROUPS	Total number of unique keys
REDUCE_INPUT_RECORDS	Total number of records read by all reduce tasks
REDUCE_OUTPUT_RECORDS	Total number of records written by all reduce tasks
REDUCE_SHUFFLE_BYTES	Total number of bytes of output from map tasks copied during shuffle to reducers
SPILLED_RECORDS	Total number of records spilled to disk by all map and reduce tasks
SPLIT_RAW_BYTES	Total number of bytes consumed for metadata (offset + length) during splits
VIRTUAL_MEMORY_BYTES	Total number of virtual memory bytes consumed by tasks (RAM + swap)



Knowledge Check	
You can use the built-in counters to validate that:	
1.	The correct number of bytes were read and written
2.	The correct number of tasks were launched and successfully ran
3.	The amount of CPU and memory consumed is appropriate for your job and cluster nodes
4.	The correct number of records were read and written
► 5.	All of the above



Learning Goals

1. Work with counters
- ▶ 2. Use the MCS to monitor jobs
3. Manage and display jobs, history, and logs
4. Write unit tests for MapReduce programs

© 2014 MapR Technologies 10

Displaying Jobs in MCS

Cluster Name: training
Logged in as mapr | Logout | Support | Help | impersonate | Log Out
Manage Journals v. 4.1.3.754.QA_Yan.MapReduce.D-A.1

Navigation: Cluster, Dashboard, Nodes, MapR Viewmap, Jobs, MapR Tables, HDFS, HDFS Metrics, HDFS Metrics, NFS HA, and more...

Dashboard - Jobs

Filter: word count

Status	Job Name	User	Start	Map	Reduce	Duration	Id	Submitted	End	Priority	Avg MAD	Max MAD
Running	word count	user01	13:45:23 11/19/2014	100%	100%	38s 244ms	_0118_0002	13:45:22 11/19/2014	13:45:51 11/19/2014	NORMAL	5s 321ms	5s 321ms

<https://webserver:8443>

© 2014 MapR Technologies 11

Getting Job Summaries in MCS

Dashboard - Jobs - word count (201411191318_0002)

word count

ID: job_201411191318_0002 Start Time: 13:45:23 11/19/2014 Duration: 38s 244ms Map Progress: 100% End Time: 13:45:51 11/19/2014 Reduce Progress: 100%

User: user01

Info Tasks

Show Map Tasks Show Reduce Tasks Show Startup/Cleanup Tasks Filter

Status	Id	Type	Primary Attempt	Start	End	Duration	Local	Node
REDUCING	_02_r_000000_0	REDUCE	_02_r_000000_0	13:45:47 11/19/2014	13:45:55 11/19/2014	7s 904ms	unknown	mapr-training
MAPPING	_02_m_000000_0	MAP	_02_m_000000_0	13:45:49 11/19/2014	13:45:49 11/19/2014	5s 310ms	remote	mapr-training

© 2014 MapR Technologies 12



 Getting Task Summaries in MCS

Dashboard | Jobs | word count 201411191318_0002 | task_201411191318_0002_r_000000

task_201411191318_0002_r_000000

Type: REDUCE
Primary Attempt Id: attempt_201411191318_0002_r_000000_0
Start Time: 13:45:47 11/19/2014
End Time: 13:45:55 11/19/2014 Duration: 7s 904ms

Task Attempts

Status	Id	Type	Progress	Start	Finish	Shuffle End	Sort End	Duration	Node	Log
OK	_r_000000_0	REDUCE	100%	13:45:47 11/19/2014	13:45:55 11/19/2014	13:45:54 11/19/2014	13:45:54 11/19/2014	7s 914ms	mapr-training	http://mapr-m.../attempt_201411191318_0002_r_000000_0

© 2014 MapR Technologies  13

 Getting Task Details in MCS

Dashboard | Jobs | word count 201411191318_0002 | attempt_201411191318_0002_r_000000_0

attempt_201411191318_0002_r_000000_0

Duration: 7s 914ms

Task Attempts (100% complete)

Map/Reduce Counters	Total
Comitted Input Records	-
Map Input Records	-
Map Output	-
Map Reducer Input	-
Map Reducer Output	-
Map Reduced Records	-
Reduce Input Count	36
Reduce Input Records	36
Reduce Output	-
Reduce Shuffle Bytes	1,380
Reduce Shuffled Records	-
Reduce Spills	-
Reduce Spilled Records	-
Reduce Total Spills	36

Map/Reduce Throughput Counters

Map	Reduce
Input Bytes Sec	-
Input Records Sec	-
Output Bytes Sec	5
Output Records Sec	5
Reduce Bytes Sec	-
Reduce Records Sec	-
shuffle Bytes Sec	1,380

File System Counters

Total	
Local Bytes Read	14,870
Local Bytes Written	14,870
Map Input Groups	2,040

MAPR 14

 Getting Log Files From MCS

Task Logs: 'attempt_201411191318_0002_r_000000_0'

student logs

student_log

student_log

student_log

```
2014-11-19 13:45:54.947 INFO mapred.Child [main]: 2014-11-19 13:45:54.947, 1239844112 pid: 21420
2014-11-19 13:45:54.944 WARN mapreduce.Counters [RPC Client (1237272949)] connection to /23.12.3.144:975 from 50b.201411191318_0002: Group org.apache.hadoop.mapreduce.TaskAttemptCounter has 0 counters: 0
2014-11-19 13:45:54.238 INFO Configuration.deprecation [main]: slave.host.name is deprecated. Instead, use dfs.datanode.hostname
2014-11-19 13:45:54.238 INFO Configuration.deprecation [main]: mapred.job.tracker is deprecated. Instead, use mapred.job.tracker.script.number.args
2014-11-19 13:45:54.138 INFO mapred.Child [main]: Starting task attempt_201411191318_0002_r_000000_0
2014-11-19 13:45:54.138 INFO mapred.Task [main]: Using configuration from /etc/hadoop-mapreduce2/conf
2014-11-19 13:45:53.135 INFO jvm.JvmMetrics [main]: Initializing JVM Metrics with processName=HDFS/2, sessionID=1395834375
2014-11-19 13:45:53.135 INFO jvm.JvmMetrics [main]: Using configuration from /etc/hadoop-mapreduce2/conf
2014-11-19 13:45:53.371 WARN util.ProcessTree [main]: /proc/gid/status does not have information about user mapr user(mapr). Can not track ownership of /proc/gid/status
2014-11-19 13:45:53.371 INFO util.ProcessTree [main]: /proc/gid/status does not have information about user mapr user(mapr). Can not track ownership of /proc/gid/status
2014-11-19 13:45:54.192 INFO mapred.Merger [main]: Merging 1 sorted segments
2014-11-19 13:45:54.192 INFO mapred.Merger [main]: Merging 1 sorted segments
2014-11-19 13:45:54.224 INFO mapred.Merger [main]: Merging 1 sorted segments
2014-11-19 13:45:54.224 INFO mapred.Merger [main]: Merging 1 sorted segments
2014-11-19 13:45:54.244 INFO mapred.Merger [main]: Merging 1 sorted segments
2014-11-19 13:45:54.245 INFO mapred.Task [main]: Task attempt_201411191318_0002_r_000000_0 is done. And is in the process of committing
2014-11-19 13:45:55.329 INFO mapred.Task [main]: Task attempt_201411191318_0002_r_000000_0 is allowed to commit now
2014-11-19 13:45:55.378 INFO mapred.Task [main]: Task attempt_201411191318_0002_r_000000_0 is committed successfully
2014-11-19 13:45:55.378 INFO mapred.Task [main]: Task attempt_201411191318_0002_r_000000_0
```

postfile.out.log

Learning Goals

1. Work with counters
2. Use the MCS to monitor jobs
- ▶ 3. Manage and display jobs, history, and logs
4. Write unit tests for MapReduce programs

© 2014 MapR Technologies 16

Tracking Status of Launched Jobs (1)

```
$ hadoop job -list
1 jobs currently running
JobID      State      StartTime   UserName  Priority  SchedulingInfo
job_201404051135_0322      1  1396917391960 user20    NORMAL    NA

$ hadoop job -status job_201404051135_0322
Job: job_201404051135_0322
file: maprfs://var/mapr/cluster/mapred/jobTracker/staging/user20/.staging/
job_201404051135_0322/job.xml
tracking URL: http://ip-10-170-165-141:50030/jdbdetails.jsp?jobid=job_201404051135_0322
map() completion: 0.024717983
reduce() completion: 0.0
Counters: 19
<output omitted>
```

© 2014 MapR Technologies 17

Tracking Status of a Launched Job (2)

© 2014 MapR Technologies 18



Stopping a Launched Job

```
$ hadoop job -list
1 jobs currently running
JobID      State    StartTime   UserName Priority SchedulingInfo
job_201404051135_03221    1396917391960    user20  NORMAL  NA

$ hadoop job -kill job_201404051135_03221
```

© 2014 MapR Technologies 19

Modifying a Job Priority

When job is launched using API

```
Configuration conf = new Configuration();
conf.set("mapred.job.priority", "VERY_HIGH");
Job job = new Job(conf, "receipts" + System.getProperty("user.name"));
```

When job is launched using a configuration file (or passing -D mapred.job.priority=X at job submission)

```
<configuration>
<property>
<name>mapred.job.priority</name><value>HIGH</value>
</property>
</configuration>
```

After job is submitted (but before it is run) using hadoop CLI

```
$ hadoop job -set-priority job_201311221648_0039 VERY_LOW
13/11/26 14:59:12 INFO fs.JobTrackerWatcher: Current Running JobTracker is:
ip-10-198-41-188/10.198.41.188:9001
Changed job priority.
```

R 20

MapR Feature: Label-Based Scheduling

© 2014 MapR Technologies 21



 Using Label-Based Scheduling

Display node labels

```
$ hadoop job -showlabels
Node labels:
Cento8002 : [big, QA, IBM]
Cento8005 : [small]
Cento8006 : [medium]
Cento8004 : [small]
Cento8001 : [big, QA, CISCO]
Cento8003 : [big, ENG, HP]
```

Submit job with a label

```
$ hadoop jar . . . -D mapred.job.label=big . . .
```

© 2014 MapR Technologies  22

 Apache Commons Logging (JCL)

© 2014 MapR Technologies  23

 Logging Information

```
private static Log log = LogFactory.getLog(MyClass.class);
. . .
public void map( . . . ) {
    if(debugcondition) {
        log.debug("debug log message");
        // OR
        System.out.println("debug output");
    } else if(errorcondition) {
        log.error("error log message");
        // OR
        System.err.println("error output");
    }
}
```

© 2014 MapR Technologies  24

Viewing the History of a Job

```
$ hadoop job -history TERASORT.0UT/
Hadoop job: job_201311211623_0015
...
Status: SUCCESS
...
|Job Counters |Aggregate execution time of mappers(ms)|0
|0 |28,978
...
Task Summary
Setup 1 1 0 0 23-Nov-2013 15:08:25 23-Nov-2013 15:08:26
(sec)
...
Analysis
Time taken by best performing map task task_201311211623_0015_m_000001: 12sec
...

```

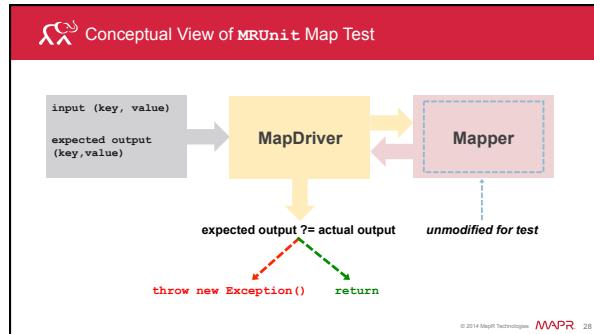
Learning Goals

1. Work with counters
2. Use the MCS to monitor jobs
3. Manage and display jobs, history, and logs
- ▶ 4. Write unit tests for MapReduce programs

Summary of MRUnit

- Testing harness for MapReduce based on JUnit
 - Developed at Cloudera (<http://mrunit.apache.org>)
- Apache*
MRUnit™
- Uses LocalJobRunner to execute code





Example MRUnit Code: Map-only (1)

```

public class ReceiptsTest {
    private static MapDriver<LongWritable, Text, Text, Text> mapDriver;

    @Before
    private static void setUp() {
        Receipts.ReceiptsMapper mapper = new Receipts.ReceiptsMapper();
        mapDriver = MapDriver.newMapDriver(mapper);
    }
}

```

© 2014 MapR Technologies 29

Example MRUnit Code:Map-only (2)

```

@Test
private static void testMapper() throws IOException,
InterruptedException {
    final LongWritable key = new LongWritable(0);
    Text value = new Text("1901 588 525 63 588 525
63 ..... ....");
    mapDriver
        .withInput(key, value)
        .withOutput(new Text("summary"), new Text ("1901_63"))
        .runTest();
}

```

© 2014 MapR Technologies 30

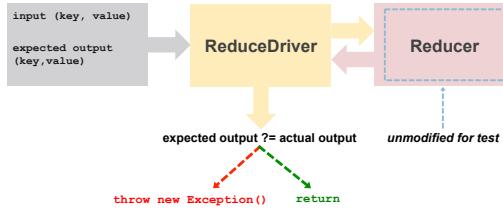


Example MRUnit Code:map-only (3)

```
public static void main(String[] args) {
    setUp();
    try { testMapper(); }
    catch (Exception e) { System.err.println("exception: " +
e.toString()); }
    return;
}
```

© 2014 MapR Technologies 31

Conceptual View of MRUnit Reducer Test



© 2014 MapR Technologies 32

Refer to Lab Guide for Exercise Instructions





 Summary

- Counters
- MCS
- Launch, manage, and monitor
- Test

© 2014 MapR Technologies. MAPR. 34

 Next Steps



© 2014 MapR Technologies. MAPR. 35



MAPR Academy
Managing Performance

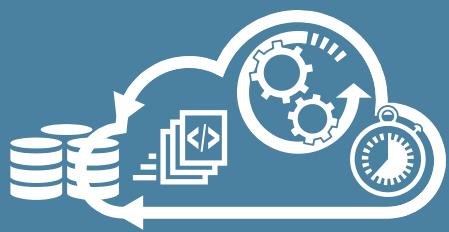
© 2014 MapR Technologies. MAPR. 1

 Learning Goals

- ▶ 1. Review components of MapReduce performance
- 2. Enhance performance in MapReduce jobs
- 3. Describe MapR performance enhancements

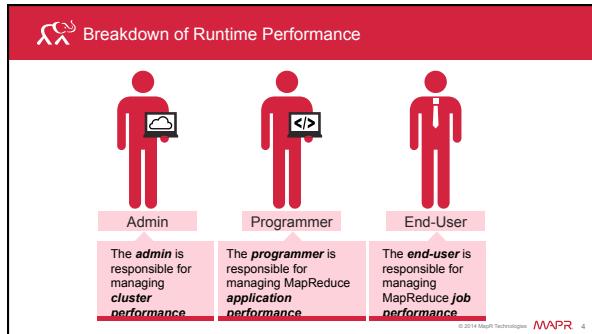
© 2014 MapR Technologies. MAPR. 2

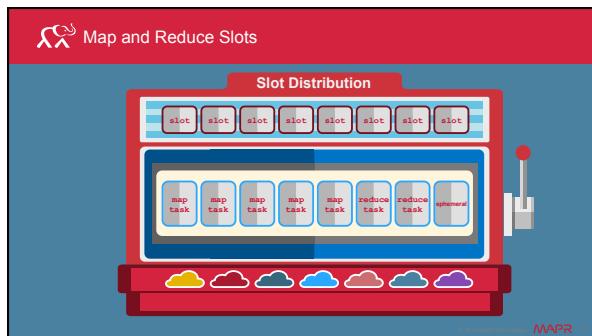
 Performance Tuning Tips



© 2014 MapR Technologies. MAPR. 3







 Knowledge Check

Slot configuration is calculated by:

1. CPU
2. Memory
3. Disk
- ▶ 4. All of the above

© 2014 MapR Technologies  6



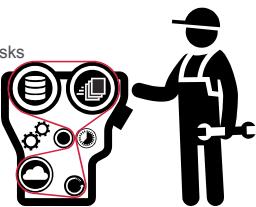
Learning Goals

1. Review components of MapReduce performance
- ▶ 2. Enhance performance in MapReduce jobs
3. Describe MapR performance enhancements

© 2014 MapR Technologies  7

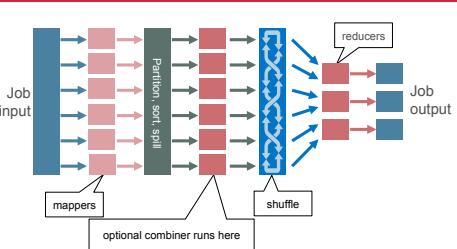
Performance Enhancements

- Using a custom combiner
- Compressing map results
- Modifying number of reduce tasks
- Using speculative execution
- Reusing a JVM
- Configuring sort properties
- Configuring Java properties



© 2014 MapR Technologies  8

Custom Combiners



© 2014 MapR Technologies  9



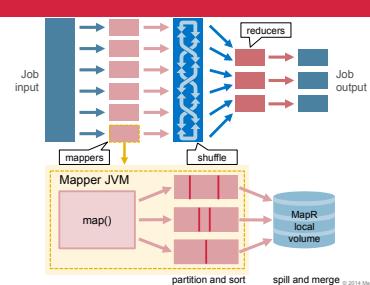
Example: Using a Combiner

```
public class UniversityDriver extends Configured implements Tool {
    public int run(String[] args) throws Exception {
        job.setCombinerClass(MyCombiner.class);
    }
}

public class MyReducer extends Reducer<Text,IntWritable,Text,FloatWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException, InterruptedException {
    }
}
```

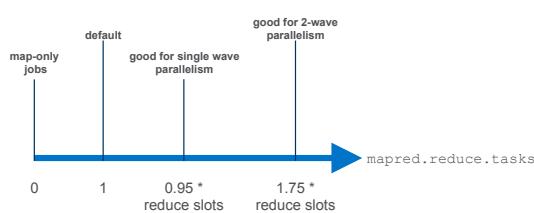
© 2014 MapR Technologies  10

Map Output Compression



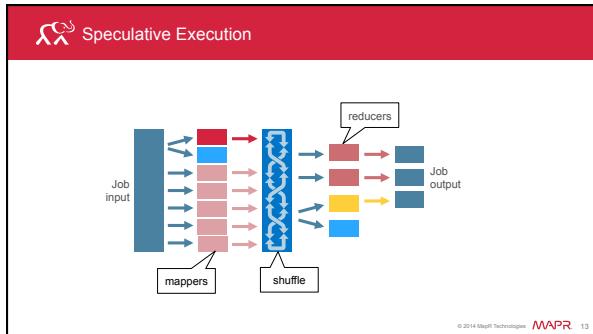
© 2014 MapR Technologies  11

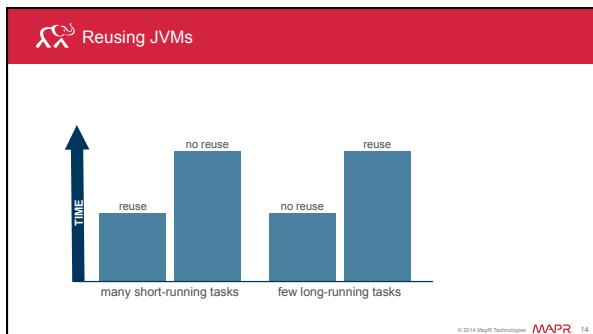
Tuning Number of Reduce Tasks

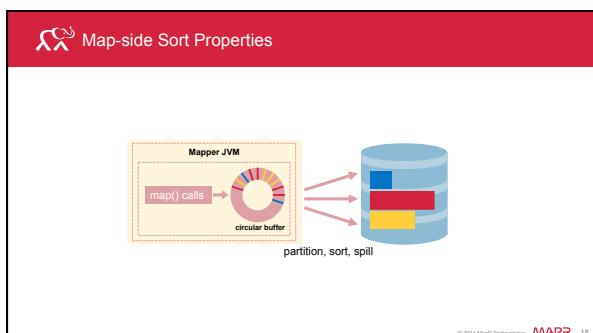


© 2014 MapR Technologies  12









 Limiting Memory Consumption

property	description
<code>mapred.map.child.java.opts</code>	JVM properties that apply to map tasks
<code>mapred.reduce.child.java.opts</code>	JVM properties that apply to reduce tasks
<code>mapred.child.java.opts</code>	JVM properties that apply to both map and reduce tasks
<code>mapred.child.ulimit</code>	Maximum size of all virtual memory consumed by a task and its children

© 2014 MapR Technologies  16

 Knowledge Check

Some performance enhancements that can be used for MapReduce jobs are:

1. Using a custom Combiner (best if the number of keys is significantly less than the number of records)
2. Compressing Mapper results
3. Modifying number of reduce tasks
- ▶ 4. All of the above

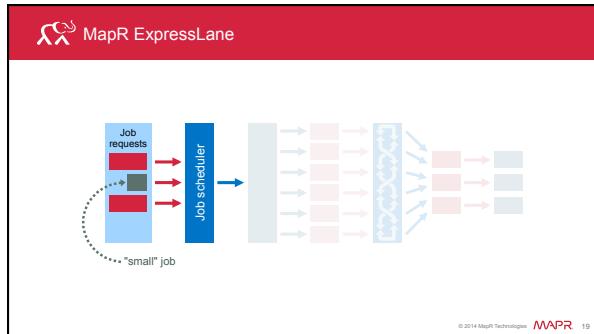
© 2014 MapR Technologies  17

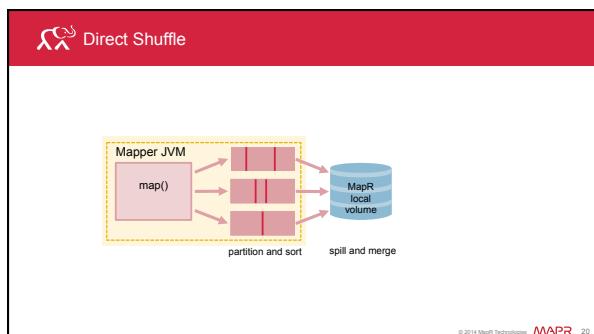
 Learning Goals

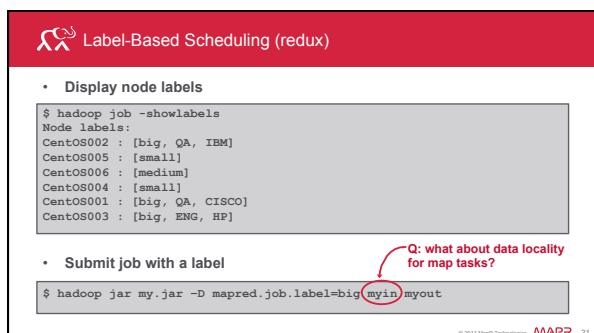
1. Review components of MapReduce performance
2. Enhance performance in MapReduce jobs
- ▶ 3. Describe MapR performance enhancements

© 2014 MapR Technologies  18











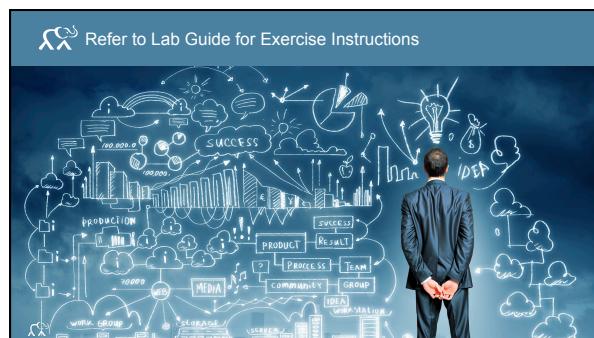
Node Labels and Topology

```
$ hadoop job -showlabels
Node labels:
CentOS001 : [small]
CentOS002 : [small]
CentOS003 : [medium]
CentOS004 : [medium]
CentOS005 : [large]
CentOS006 : [large]
```

Hostname	Physical IP(s)	File...	Physical Topology
CentOS001	10.10.82.51	0 ago	/row1/rack1/CentOS001
CentOS002	10.10.82.52	0 ago	/row1/rack1/CentOS002
CentOS003	10.10.82.54	0 ago	/row1/rack1/CentOS003
CentOS004	10.10.82.55	0 ago	/row1/rack1/CentOS004
CentOS005	10.10.82.59	0 ago	/row1/rack1/CentOS005
CentOS006	10.10.82.60	0 ago	/row1/rack1/CentOS006
	10.10.82.61	0 ago	
	10.10.82.62	0 ago	

same label ↔ same topology

© 2014 MapR Technologies 22



Summary

- Components of MapReduce performance
- Enhancing performance of MapReduce jobs

© 2014 MapR Technologies 24





MAPR Academy

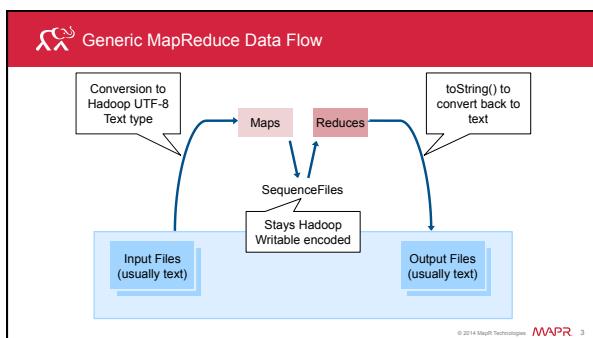
Working with Data

© 2014 MapR Technologies. MAPR 1

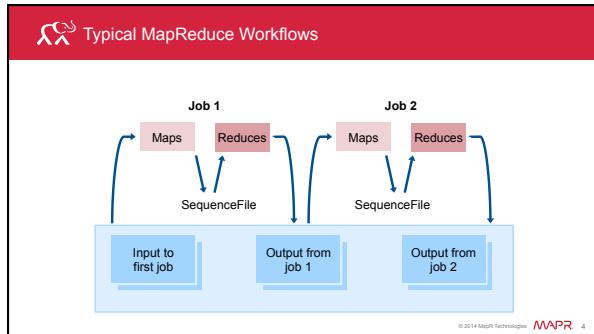
 Learning Goals

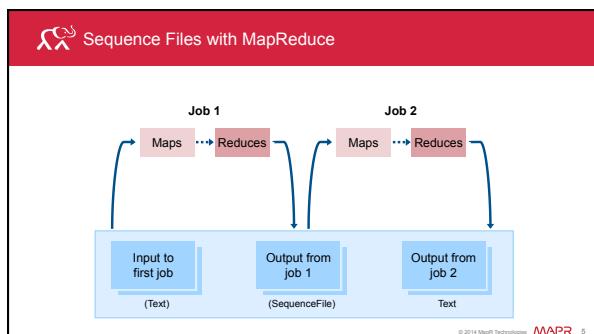
- ▶ 1. Working with sequence files
- 2. Working with the distributed cache
- 3. Working with HBase

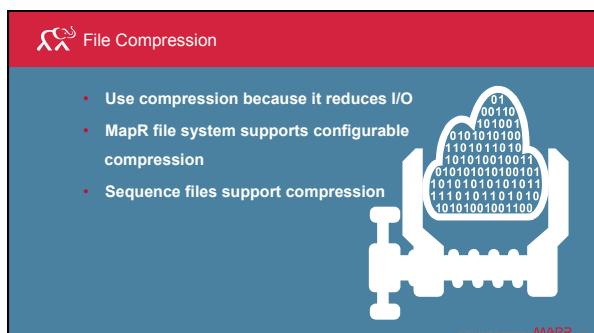
© 2014 MapR Technologies. MAPR 2







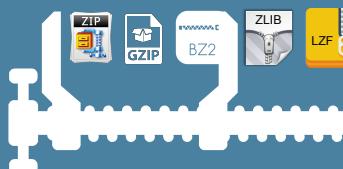






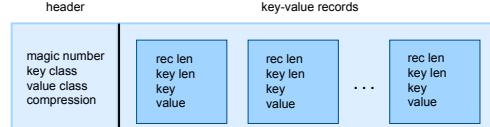
Compressing Sequence Files

- NONE → do not compress
- RECORD → compress values only (not keys)
- BLOCK → compress sequences of records in blocks



© 2014 MapR Technologies. **MAPR** 8

Internal Layout of a Sequence File



header

magic number
key class
value class
compression

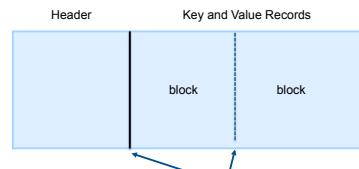
key-value records

rec len
key len
key
value

...
rec len
key len
key
value

© 2014 MapR Technologies. **MAPR** 9

Sync Markers



Header

Key and Value Records

block

block

Markers (tunable ~ %64k)

© 2014 MapR Technologies. **MAPR** 9



MapReduce SequenceFile Syntax

```
job.setOutputFormatClass(SequenceFileOutputFormat.class);
SequenceFileOutputFormat.setCompressOutput(job, true);
SequenceFileOutputFormat.setOutputCompressorClass(job, GzipCodec.class);
SequenceFileOutputFormat.setOutputCompressionType(job,
CompressionType.BLOCK);
job.setInputFormatClass(SequenceFileInputFormat.class);
```

© 2014 MapR Technologies  10

Knowledge Check

- Which of the following are true of sequence files?
1. Intermediate data from the Mapper is stored as sequence files
 2. Sequence files can be compressed
 3. Using sequence files can save space and time in complex workflows comprising multiple MapReduce jobs
 - ▶ 4. All of the above
 5. 1 & 3

© 2014 MapR Technologies  11

Learning Goals

1. Working with sequence files
- ▶ 2. Working with the distributed cache
3. Working with HBase

© 2014 MapR Technologies  12





Populate the Distributed Cache Using the hadoop Command

- Pass plain text files to Mapper and Reducer classes

```
hadoop jar . . . --files file1,file2 . . .
```

- Pass Java jar files to Mapper and/or Reducer classes

```
hadoop jar . . . --libjars jar1,jar2 . . .
```

- Pass JAR/TAR or ZIP/GZIP archives to Mapper/Reducer classes

```
hadoop jar . . . --archives archive1,archive2 . . .
```

© 2014 MapR Technologies  13



Populate the Distributed Cache using the API

```
// driver class
DistributedCache.addCacheFile(new Path("/user/user01/myfile").toUri(),
getConf());
DistributedCache.addCacheArchive(new Path("/user/user01/
myarchive.zip").toUri(), getConf());
```

© 2014 MapR Technologies  14



Describe How to Consume the Distributed Cache

```
// Mapper or Reducer class
. . .
Path[] localFiles, localArchives;

public void setup(Context context) throws IOException {
    Configuration conf = context.getConfiguration();
    localFiles = DistributedCache.getLocalCacheFiles(conf);
    localArchives = DistributedCache.getLocalArchives(conf);
}

// consume cache in map or reduce method
```

© 2014 MapR Technologies  15

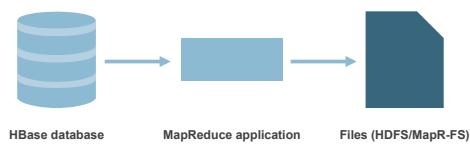


Learning Goals

1. Working with sequence files
2. Working with the distributed cache
- ▶ 3. Working with HBase

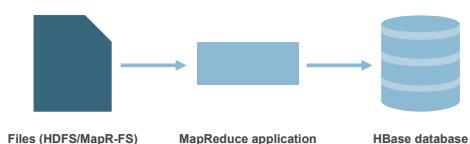
© 2014 MapR Technologies **MAPR** 16

Using HBase as a Source



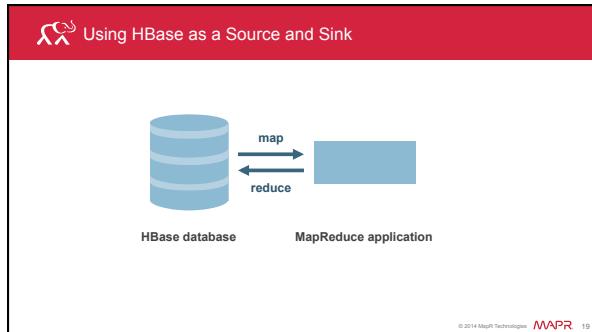
© 2014 MapR Technologies **MAPR** 17

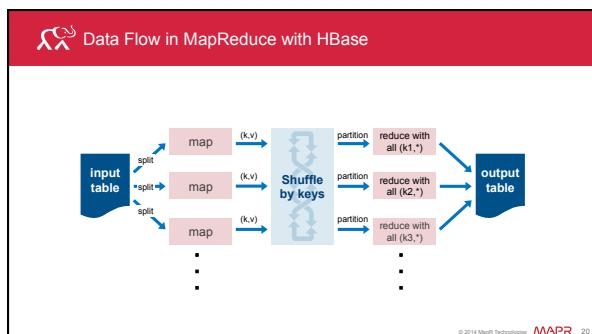
Using HBase as a Sink

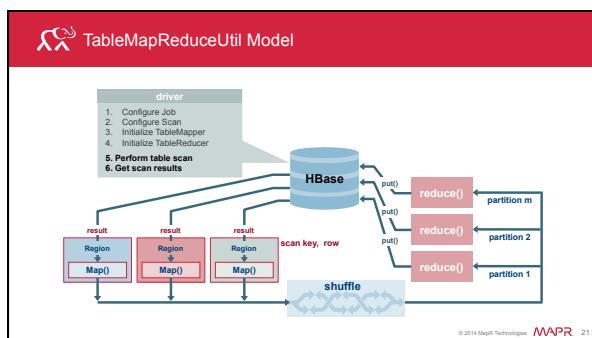


© 2014 MapR Technologies **MAPR** 18











Use a Flat-Wide Schema

```
create '/mapr/my.cluster.com/user/user01/trades_flat',
{NAME=>'price', VERSIONS=>1000000},
{NAME=>'vol', VERSIONS=>1000000},
{NAME=>'stats'}
```

rowkey	price:00	price:10	...	price:23	vol	stats:min	stats:max	stats:mean
AMZN_20131010								
GOOG_20131010								
CSCO_20131010								
...								

© 2014 MapR Technologies  22

Driver Example

```
public int run(String[] args) throws Exception {
    // set up the job
    Job job = Job.getInstance(getConf(), getClass().getSimpleName());
    job.setJarByClass(getClass());
    //Create and configure a Scan instance.
    Scan scan = new Scan();
    scan.setMaxVersions();
    scan.addFamily(Bytes.toBytes("price"));
    // initialize the mapper and reducer
    TableMapReduceUtil.initTableMapperJob(INPUT_TABLE, scan,
        StockMapper.class, Text.class, LongWritable.class, job);
    TableMapReduceUtil.initTableReducerJob(
        OUTPUT_TABLE, StockReducer.class, job);
    return job.waitForCompletion(true) ? 0 : 1;
}
```

© 2014 MapR Technologies  23

TableMapper Example

```
public class StockMapper extends TableMapper<Text, LongWritable> {
    @Override
    protected void map(ImmutableBytesWritable key, Result row,
        Context context) {
        String keyString = Bytes.toString(key.get());
        final Text outKeyText = new Text(keyString);
        // for all columns in input row
        for (KeyValue col : row.list()) {
            byte[] cellValueBytes = col.getValue();
            long price = Bytes.toLong(cellValueBytes);
            context.write(outKeyText, new LongWritable(price));
        }
    }
}
```

© 2014 MapR Technologies  24



 TableReducer Example

```

public class StockReducer extends
    TableReducer<Text, LongWritable, ImmutableBytesWritable> {
    @Override
    protected void reduce(Text key, Iterable<LongWritable> prices,
        Context context) {
        long min=Long.MAX_VALUE;
        long temp = 0L;
        for (LongWritable price : prices) {
            temp=price.get();
            if(temp < min) min=temp;
        }
        Put put = new Put(Bytes.toBytes(key.toString()));
        put.add(OUTPUT_FAMILY, OUTPUT_COLUMN,
            Bytes.toBytes(Float.toString((float)min/100)));
        context.write(key, put);
    }
}

```

 Refer to Lab Guide for Exercise Instructions



 Summary

1. Working with sequence files
2. Working with the distributed cache
3. Working with HBase

© 2014 MapR Technologies  27





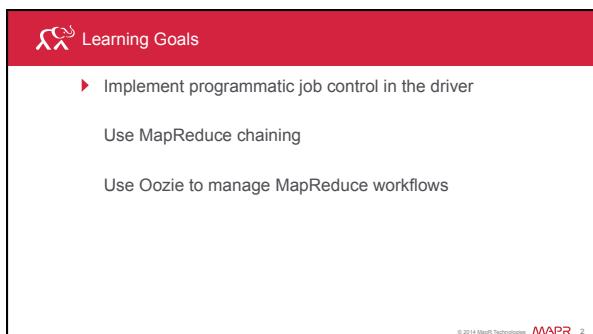




MAPR® Academy

Launching Jobs

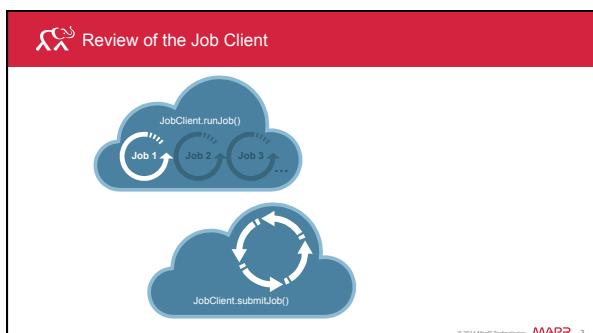
© 2014 MapR Technologies. **MAPR**. 1



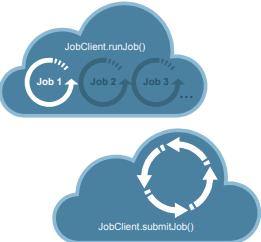
 **Learning Goals**

- ▶ Implement programmatic job control in the driver
- Use MapReduce chaining
- Use Oozie to manage MapReduce workflows

© 2014 MapR Technologies. **MAPR**. 2

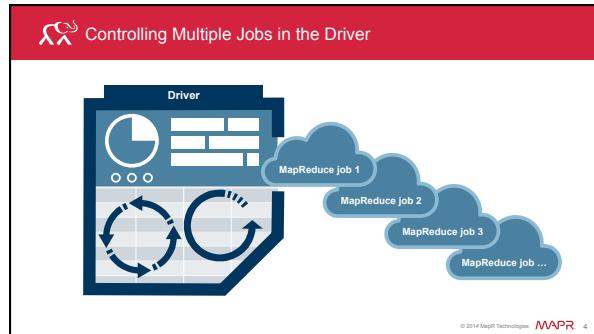


 **Review of the Job Client**



© 2014 MapR Technologies. **MAPR**. 3





 Submitting Two Jobs

```
public class TwoJobs {
    public static class Map extends MapReduceBase implements Mapper<...> {
        public void map() { ... }
    }
    public static class Reduce extends MapReduceBase implements Reducer<...> {
        public void reduce() { ... }
    }
    public static class Map2 extends MapReduceBase implements Mapper<...> {
        public void map() { ... }
    }
    public static class Reduce2 extends MapReduceBase implements Reducer<...> {
        public void reduce() { ... }
    }
}
```

© 2014 MapR Technologies  5

 Submitting Two Jobs (2)

```
public static void main(String[] args) throws Exception {
    JobConf conf1 = new JobConf(TwoJobs.class);
    conf1.setMapperClass(Map.class);
    conf1.setReducerClass(Reduce.class);
    JobClient.runJob(conf1);
    JobConf conf2 = new JobConf(TwoJobs.class);
    conf2.setMapperClass(Map2.class);
    conf2.setReducerClass(Reduce2.class);
    JobClient.runJob(conf2);
}
```

© 2014 MapR Technologies  6



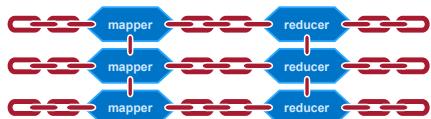
Learning Goals

- Implement programmatic job control in the driver
- ▶ Use MapReduce chaining
- Use Oozie to manage MapReduce workflows

© 2014 MapR Technologies  7

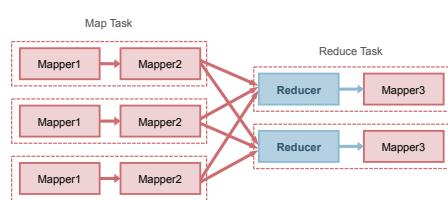
Job Chaining

- ChainMapper and ChainReducer "chain" successive mappers or successive reducers
- Data transfer is in memory, rather than disk



© 2014 MapR Technologies  8

MapReduce Chaining



© 2014 MapR Technologies  9



Job Chaining Pattern

```
...
JobConf conf = new JobConf(true);
...
JobConf mapAConf = new JobConf(false);
ChainMapper.addMapper(conf, AMap.class, LongWritable.class, Text.class,
Text.class, Text.class, true, mapAConf);

JobConf mapBConf = new JobConf(false);
ChainMapper.addMapper(conf, BMap.class, Text.class, Text.class,
LongWritable.class, Text.class, false, mapBConf);

JobConf reduceConf = new JobConf(false);
ChainReducer.setReducer(conf, Reduce.class, LongWritable.class, Text.class,
Text.class, IntWritable.class, true, reduceConf);

JobClient.runJob(conf);
```

© 2014 MapR Technologies 10

Knowledge Check

Which of the following are true of MapReduce chaining

1. Can be used to manage multiple MapReduce jobs
2. Data from mappers are passed onto next mappers or reducers in memory
3. You must use the Mapred package for chaining
- ▶ 4. All of the above
5. 2 & 3 only

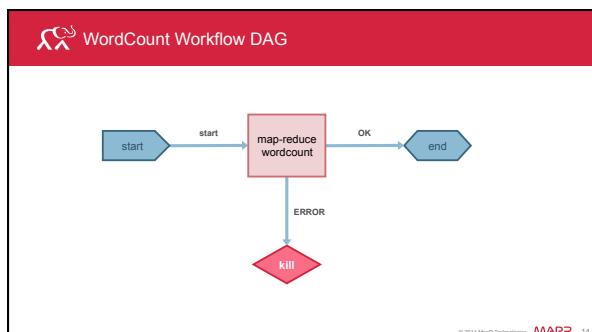
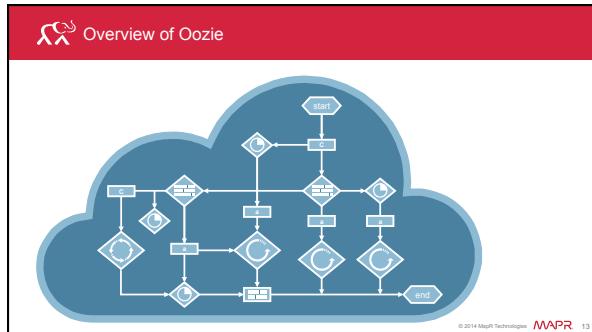
© 2014 MapR Technologies 11

Learning Goals

- Implement programmatic job control in the driver
- Use MapReduce chaining
- ▶ Use Oozie to manage MapReduce workflows

© 2014 MapR Technologies 12

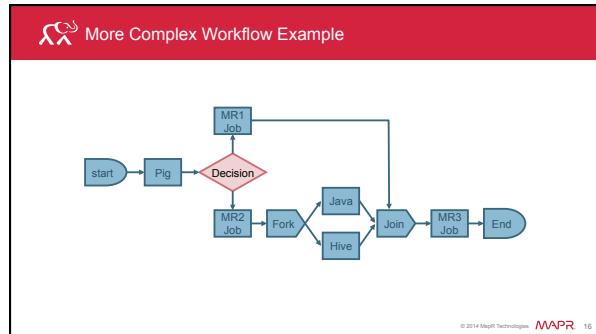




 Oozie Workflow XML

```
<workflow-app name="wordcount-wf" xmlns="uri:oozie:workflow:0.1">
  <start to="wordcount"/>
  <action name="wordcount">
    <configuration>
      <job-tracker${jobTracker}/>
      <name-node${nameNode}/>
      <configuration>
        <property>
          <name>mapred.mapper.class</name>
          <value>org.apache.hadoop.mapred.MapReduceMapper</value>
        </property>
        <property>
          <name>mapred.reducer.class</name>
          <value>org.apache.hadoop.mapred.ReduceTaskInputOutputProcessor$IdentityReducer</value>
        </property>
        <property>
          <name>mapred.input.dir</name>
          <value>${inputDir}</value>
        </property>
        <property>
          <name>mapred.output.dir</name>
          <value>${outputDir}</value>
        </property>
      </configuration>
    </configuration>
    <ok to="end"/>
    <error to="end"/>
  </action>
  <kill name="kill1">
    <message>Something went wrong: ${wf:errorCode('wordcount')}</message>
  </kill1>
  <end name="end"/>
</workflow-app>
```



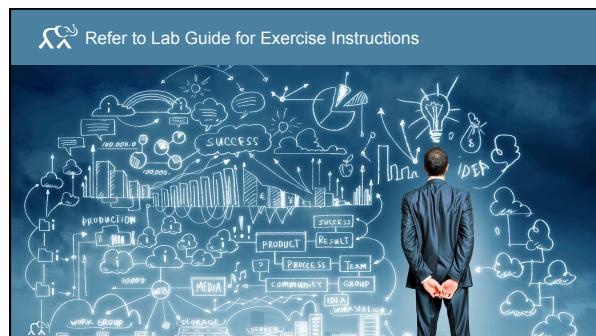


Knowledge Check

Apache Oozie is a client-server workflow engine for MapReduce and the ecosystem. Which of the following are true of Oozie?

- ▶ 1. Oozie workflows can be modified by editing the corresponding XML file
- 2. Workflows are represented as cyclic graphs
- 3. Oozie workflows handle linear flow
- 4. All of the above

© 2014 MapR Technologies 17



Summary

- Launching multiple MapReduce jobs
- MapReduce Chaining
- Apache Oozie

© 2014 MapR Technologies. MAPR. 19

Next Steps



© 2014 MapR Technologies. MAPR. 20



MapR Academy

Using non-Java programs (Streaming MapReduce)

© 2014 MapR Technologies.  1

Learning Goals

- ▶ 1. Overview using non-Java programming (streaming)
- 2. Define the programming contract for mappers & reducers
- 3. Monitoring and debugging MapReduce streaming jobs

© 2014 MapR Technologies.  2

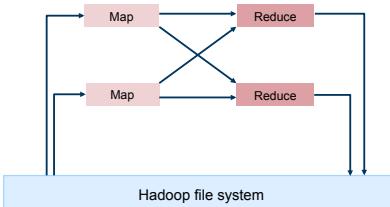
MapReduce Streaming

- MapReduce streaming enables using other languages
- Streaming *!= pipes* (C++ equivalent)
- Streaming *may* introduce some performance penalty
- Streaming *may* improve performance

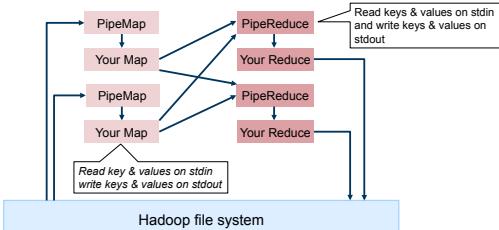
© 2014 MapR Technologies.  3



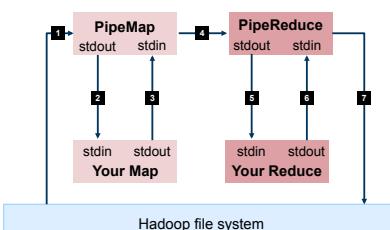
Review the Typical MapReduce Job


© 2014 MapR Technologies. 4

MapReduce Streaming


© 2014 MapR Technologies. 5

Data Flow in a Streaming Job


© 2014 MapR Technologies. 6



Linux Commands in Streaming (1)

```
# Read input from local filesystem
# Write output to home dir on cluster
• # eg. /mapr/my.cluster.com/user/...
export CONTRIB=/opt/mapr/hadoop/hadoop*/contrib/streaming
export STREAMINGJAR=hadoop-*.streaming.jar
export THEJARFILE=$CONTRIB/$STREAMINGJAR

hadoop jar $THEJARFILE \
  

    -input file:///etc/passwd  

    -mapper '/bin/cat' \\  

    -reducer '/bin/cat'  

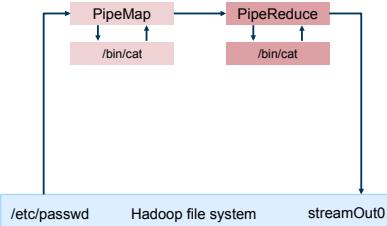
  


```

© 2014 MapR Technologies  7

Linux Commands in Streaming (2)



© 2014 MapR Technologies  8

Linux Commands in Streaming (3)

```
more /etc/passwd
root:x:0:0:root:/root:/bin/bash
root:x:1:1:root:/root:/bin/sh
bin:x:2:2:bin:/bin:/bin/sh
sync:x:3:3:sync:/dev:/bin/sh
games:x:4:65534:sync:/bin:/bin/sync
games:x:5:60:games:/usr/games:/bin/sh
man:x:6:12:man:/var/cache/man:/bin/sh
...
more /mapr/my.cluster.com/user/user01/streamOut0/
part-0000
avahi-autopid:x:103:110:Avahi autopid daemon,,,:/var/lib/avahi-autopid:/bin/false
avahi:x:104:111:Avahi mDNS daemon,,,:/var/run/avahi-daemon:/bin/false
backup:x:34:34:backup:/var/backups:/bin/sh
bin:x:2:2:bin:/bin:/bin/sh
couchdb:x:105:113:CouchDB Administrator,,,:/var/lib/couchdb:/bin/bash
daemon:x:1:1:daemon:/usr/sbin:/bin/sh
games:x:5:60:games:/usr/games:/bin/sh
...

```

© 2014 MapR Technologies  9





Knowledge Check

- What advantages does MapReduce streaming provide?
1. You can use streaming for rapid prototyping using sed/awk.
 2. It enables you to use languages other than Java (Perl, Python) to write MapReduce programs
 3. It definitely enhances performance
 4. All of the above
 - ▶ 5. 1 & 2

© 2014 MapR Technologies.  10



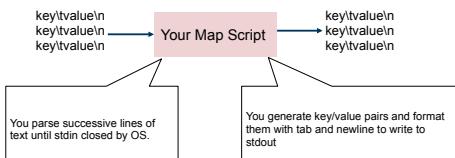
Learning Goals

1. Overview using non-Java programming (streaming)
- ▶ 2. Define the programming contract for mappers & reducers
3. Monitoring and debugging MapReduce streaming jobs

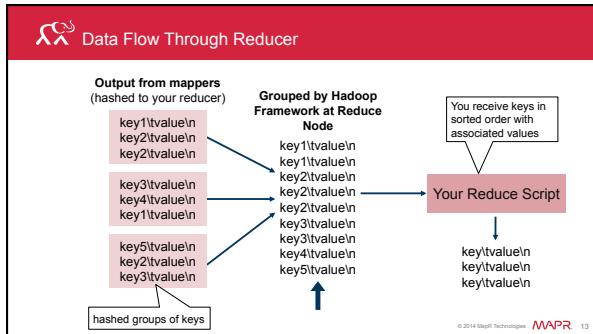
© 2014 MapR Technologies.  11

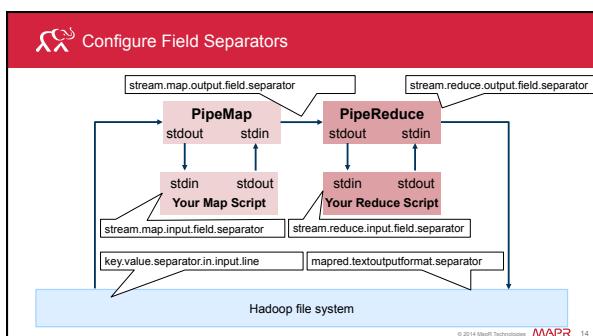


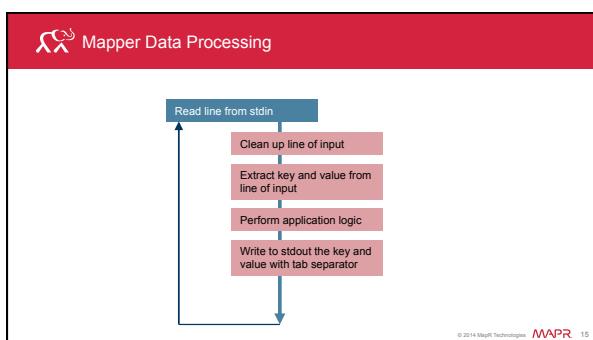
Data Flow Through Mapper


© 2014 MapR Technologies.  12







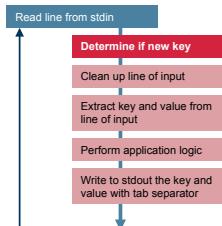


Perl Example: Mapper

```
#!/usr/bin/env perl
while (<>) { # read stdin
    chomp; # remove last char if newline from the implicit variable $_
    my ( $key,$value ) = split(/\t/,$_); # extract key and value
    print "key:".$key."\t".value."\n"; # write to stdout
}
```

© 2014 MapR Technologies  16

Reducer Data Processing



© 2014 MapR Technologies  17

Reducer Data Processing

```
#!/usr/bin/env python
import sys
currentKey = None
groupTotal = 0
for line in sys.stdin:
    line = line.strip() # remove leading and trailing whitespace
    if not line: # ignore empty lines
        continue
    keyFromLine, valueFromLine = line.split("\t", 1)
    if currentKey == keyFromLine: # key hasn't changed
        groupTotal += int(valueFromLine)
    else: # we just got new key
        if currentKey: # first line of processing
            print("%s\t%d" % (currentKey, groupTotal))
        groupTotal = int(valueFromLine) # total is just our current valueFromLine
        currentKey = keyFromLine # set the new key as our currentKey
    # The last key must be processed
if currentKey == keyFromLine
    print("%s\t%d" % (currentKey, groupTotal))
```

© 2014 MapR Technologies  18



Python MapReduce Job

```
export CONTRIB=/opt/mapr/hadoop/hadoop*/contrib/streaming
export STREAMINGJAR=hadoop-*~streaming.jar
export THEJARFILE=$CONTRIB/$STREAMINGJAR
rm -rf /mapr/my_cluster.com/user/user01/pyNoOpMapRed 2>>/dev/null
hadoop jar $THEJARFILE \
-D mapred.reduce.tasks=2 \
-mapper "python noOpMap.py" \
-file noOpMap.py \
-reducer "python noOpReduce.py" \
-file noOpReduce.py \
-input file:///$/goodInput.txt \
-output "pyNoOpMapRed"
```

© 2014 MapR Technologies 19

Launching Perl Streaming Code

```
#!/usr/bin/env bash
export CONTRIB=/opt/mapr/hadoop/hadoop*/contrib/streaming
export STREAMINGJAR=hadoop-*~streaming.jar
export THEJARFILE=$CONTRIB/$STREAMINGJAR

hadoop jar $THEJARFILE \
-mapper 'perl map.pl' \
-reducer 'perl reduce.pl' \
-file map.pl \
-file reduce.pl \
-input file:///etc/passwd \
-output perlout
```

© 2014 MapR Technologies 20

Learning Goals

1. Overview using non-Java programming (streaming)
2. Define the programming contract for mappers & reducers
3. Monitoring and debugging MapReduce streaming jobs

© 2014 MapR Technologies 21



 General Debugging Tips

- Launch map or reduce script "**standalone**" **with stdin**
- Test with **bad data**
- Run a **map-only** job to test mapper
- Run **reduce-only** with an "identity mapper" that you create
- Use **counters and status**



 Perl Mapper Updating a Counter

Perl

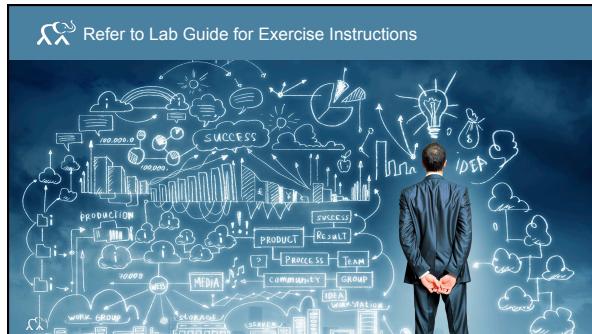
```
...  
print STDERR "reporter:counter: MyGroup,NumberOfRecords,100\n";
```

Python

```
...  
sys.stderr.write("reporter:counter: Mygroup, NumberOfRecords,100\n")
```

 Perl Reducer Updating Status

```
#!/usr/bin/env perl  
my $numRecords = 0;  
while (<>) {          # read from stdin  
    my ($key,$value) = split(/\t/);      # split input on newline  
    $numRecords += 1;                      # count num of records  
}  
print "$numRecords\n";           # after all stdin processed  
                                # write the number of records  
print STDERR "reporter:status:A Reducer Exiting\n";
```





Summary

- Non-Java programming
- Configure parameters
- Monitor and debug

© 2014 MapR Technologies. **MAPR** 26



Next Steps



© 2014 MapR Technologies. **MAPR** 27

